

# Network Layer

## Study-Ready Notes

Compiled by Andrew Photinakis

November 17, 2025

## Contents

<b>1</b>	<b>5.1 Introduction</b>	<b>2</b>
<b>2</b>	<b>5.2 Routing Algorithms</b>	<b>2</b>
2.1	Concept Overview . . . . .	2
2.2	Graph Abstraction: Modeling the Network . . . . .	2
2.3	Taxonomy of Routing Algorithms . . . . .	2
2.4	5.2.1 Link State Routing Algorithm . . . . .	3
2.5	5.2.2 Distance Vector Routing Algorithm . . . . .	4
2.5.1	Bellman-Ford Algorithm . . . . .	5
2.5.2	Comparison of Link-State vs Distance-Vector Algorithms . . . . .	6
2.5.3	5.2 Section-Wide Summary . . . . .	6
<b>3</b>	<b>5.3 Intra-AS Routing in the Internet: OSPF</b>	<b>8</b>
3.1	Need for Routing Hierarchy: Scale and Autonomy . . . . .	8
3.2	Solution: Autonomous Systems (ASs) . . . . .	9
3.3	OSPF: Canonical Intra-AS Protocol . . . . .	9
3.4	OSPF Link-State Broadcasts . . . . .	10
3.5	Advanced Features in OSPFv2 . . . . .	10
3.6	Section-Wide Summary . . . . .	11
<b>4</b>	<b>5.4 Routing Among the ISPs: BGP</b>	<b>13</b>
4.1	5.4.1 Role of BGP . . . . .	13
4.2	5.4.2 Advertising BGP Route Information . . . . .	14
4.3	5.4.3 Determining the Best Routes . . . . .	15
4.4	5.4.4 IP-Anycast . . . . .	18
4.5	5.4.5 Routing Policy . . . . .	18
4.6	5.4.6 Putting the Pieces Together: Obtaining Internet Presence . . . . .	19
4.7	5.4 Section-Wide Summary . . . . .	20
<b>5</b>	<b>5.5 The SDN Control Plane</b>	<b>22</b>

<b>6</b>	<b>The Four Key Characteristics of an SDN Architecture</b>	<b>22</b>
6.1	Flow-based Forwarding . . . . .	22
6.2	Separation of Data Plane and Control Plane . . . . .	22
6.3	Network Control Functions are External to Data-Plane Switches . . . . .	22
6.4	A Programmable Network . . . . .	23
<b>7</b>	<b>5.5.1 The SDN Control Plane: Controller and Applications</b>	<b>23</b>
7.1	The Communication Layer (Southbound Interface) . . . . .	24
7.2	The Network-Wide State-Management Layer . . . . .	24
7.3	The Interface to Applications (Northbound Interface) . . . . .	24
7.4	Practical Implementation: ODL and ONOS . . . . .	24
<b>8</b>	<b>5.5.2 OpenFlow Protocol</b>	<b>24</b>
8.1	Concept Overview . . . . .	24
8.2	Technical Mechanism: Key Message Types . . . . .	24
8.3	Practical View: Google's B4 Network . . . . .	25
<b>9</b>	<b>5.5.3 Data and Control Plane Interaction: An Example</b>	<b>25</b>
<b>10</b>	<b>5.5.4 SDN: Past and Future</b>	<b>25</b>
10.1	Past: The Origins of SDN . . . . .	25
10.2	Future: Network Functions Virtualization (NFV) . . . . .	25
10.3	Future: Inter-AS SDN . . . . .	25
<b>11</b>	<b>5.5 Section-Wide Summary</b>	<b>25</b>
11.1	Key Terms and Definitions . . . . .	25
11.2	Core Relationships . . . . .	25
11.3	Key Insights / Takeaways . . . . .	26

## Chapter 5: The Network Layer — Control Plane

### 1 5.1 Introduction

### 2 5.2 Routing Algorithms

#### 2.1 Concept Overview

1. Logic of how packets forwarded is done in control plane, implemented by routing algorithms
2. A good path is synonymous with "least-cost", such path populates forwarding table in each router

#### 2.2 Graph Abstraction: Modeling the Network

1. Assume graph  $G = (N, E)$
2.  $N$  = nodes in graph that represent routers in network, are points where packet-forwarding decisions are made
3.  $E$  = edges, represent physical links that connect routers
4. Costs  $c(x, y) =$  each edge  $(x, y)$  is assigned a numerical cost
  - (a) Cost is typically set by network admin, is a piece of policy
  - (b) 3 types of costs
    - i. Physical distance, a trans-oceanic link has a higher cost than a short link
    - ii. Link speed, a 1 Gbps link might have a cost of 10, while a 100 Gbps link has a cost of 1. Often inversely proportional to bandwidth
    - iii. Monetary cost, leasing a link from another provider may be more "costly"

#### 2.3 Taxonomy of Routing Algorithms

1. Centralized
  - (a) Algo computes least-cost path using complete, global knowledge of network
  - (b) Has complete map of all nodes and all edge costs
  - (c) Can be done at a single site (like an SDN controller) or replicated on every router, but key is that the input is the global network state
  - (d) Approach is known as a Link-State algorithm
2. Decentralized
  - (a) Calculation distributed among routers

- (b) Each node begins with only local knowledge (cost of its own directly attached links)
  - (c) Iterative process of communication only with neighbors, a node gradually learns least-cost paths to other nodes
  - (d) Approach known as Distance-Vector algorithm
3. Static vs Dynamic
- (a) Static: routes are changed very slowly, often by human admin manually editing costs or forwarding tables
  - (b) Dynamic: routes change automatically as network topology or link costs change.
4. Load-Sensitive vs Load-Insensitive
- (a) Load-Sensitive: link costs dynamically change to reflect current level of congestion on link. Idea to route around congestion
  - (b) Load-Insensitive: link costs are fixed and do not reflect current load

## 2.4 5.2.1 Link State Routing Algorithm

1. Centralized approach that uses complete, global knowledge of network and link costs to compute least-cost paths
2. Each node in network responsible for a simple task
  - (a) Discover neighbors and link costs
  - (b) Broadcast this info to all other nodes in network
  - (c) Receive link-state packets from all other nodes and build a complete, identical map of entire network
  - (d) Compute least-cost paths from itself to all other nodes using map
3. Dijkstra's Algorithm
  - (a) Iterative algo, finds least-cost paths in increasing order of cost
  - (b) After k-th iteration, algo has found least-cost path to nodes that are closest to source
  - (c) Since each router computes its own table independently based on received global info, a malfunctioning router can only broadcast an incorrect cost for its attached links, and cannot directly corrupt entire network's routing knowledge
  - (d) Variables
    - i. N: set of nodes to which least-cost path is definitively known
    - ii.  $D(u)$ : current cost of least-cost path from source (u) to node v. value updated as new, shorter paths found
    - iii.  $p(u)$ : predecessor node neighbor of v along current least-cost path from source. Used to build forwarding table

```

        // Initialization
1. N' = {u} // The source node is the only member
2. for all nodes v
3.     if v is a neighbor of u
4.         then D(v) = c(u,v)
5.     else D(v) = infinity
6.
// Loop
7. loop
8.     find w not in N' such that D(w) is a minimum
9.     add w to N'
10.    update D(v) for each neighbor v of w and not in N':
11.        // Does the path through w to v offer a shorter
12.        D(v) = min( D(v), D(w) + c(w,v) )
13.        // If D(v) was updated, set p(v) = w
14. until N' = N

```

(e) Pathological Oscillations

- i. Problem arises if link costs are load-sensitive, that is if cost of a link depends on amount of traffic it is carrying
- ii. Cost of link equals traffic it carries
- iii. Solution is to not use load sensitive costs and ensure routers don't run algorithm at same time. Done by randomizing when link state updates are sent or when algorithm is run

## 2.5 5.2.2 Distance Vector Routing Algorithm

1. Distance vector algorithm is decentralized approach
2. Code is it "tell your neighbors what you know about the world"
3. Algorithm is:
  - (a) Distributed, each node communicates only with direct neighbors
  - (b) Iterative, process repeats until no more information is exchanged and calculations stabilized
  - (c) Async, nodes don't need to operate in lock step
  - (d) Self-Terminating, algorithm stops when calculations converge
4. In DV, a node knows cost to its neighbors, and receives its "distance vectors" from its neighbors. A distance vector is a list of cost estimates from that neighbor to all other nodes in network

### 2.5.1 Bellman-Ford Algorithm

#### 1. Definitions of variables:

- $d_x(y)$ : Node  $x$ 's estimate of the least-cost path from  $x$  to destination  $y$ .
- $c(x, v)$ : The cost of the direct link from node  $x$  to neighbor  $v$ .
- $N(x)$ : The set of neighbors of node  $x$ .
- $d_v(y)$ : Node  $v$ 's estimate of the least-cost path from  $v$  to destination  $y$ .

#### 2. Bellman–Ford update formula:

$$d_x(y) = \min_{v \in N(x)} \{c(x, v) + d_v(y)\}$$

- Node  $x$  considers every neighbor  $v \in N(x)$ .
- For each neighbor, it computes:

$$c(x, v) + d_v(y)$$

- It chooses the neighbor  $v$  that yields the smallest total cost.

#### 3. Intuition:

The cost from node  $x$  to destination  $y$  is the minimum, over all neighbors  $v$ , of the cost to reach  $v$  plus the cost (as advertised by  $v$ ) for  $v$  to reach  $y$ .

#### 4. Distance Vector at each node:

Each node  $x$  maintains a distance vector

$$D_x = [D_x(y) : y \in \mathcal{N}]$$

representing its current estimates of  $d_x(y)$  for all destinations  $y$ .

#### 5. Update behavior:

When node  $x$  receives a distance vector  $D_v$  from neighbor  $v$ , it recomputes its own distance vector using the Bellman–Ford formula. If any entry changes, node  $x$  sends its updated distance vector  $D_x$  to all neighbors.

#### 6. Distance Vector Algorithm (Pseudo-Code):

At each node  $x$ :

1. Initialization:
2.     for all destinations  $y$  in  $\mathcal{N}$ :
3.          $D_x(y) = c(x, y)$      // If  $y$  not a direct neighbor:  $c(x, y) = \text{infinity}$
4.     for each neighbor  $w$ :
5.         send distance vector  $D_x$  to  $w$
  
6. loop
7.     wait until a distance vector  $D_v$  is received from neighbor  $v$

```

8.      for each destination y in N:
9.          // Bellman-Ford update
10.         D_x(y) = min over all neighbors u { c(x,u) + D_u(y) }

11.     if any D_x(y) changed:
12.         send updated distance vector D_x to all neighbors

13. forever

```

## 7. Pathological Dynamics

- (a) Complexity of DV is when link costs change
- (b) Cases to consider:
  - i. Link cost decreases
  - ii. Link cost increases at & count to infinity problem
- (c) Solution
  - i. Poisoned Reverse
  - ii. A simple fix for 2-node loops
    - A. if z routes to x through y, z will "lie" to y and advertise its distance to x to infinity
  - iii.

### 2.5.2 Comparison of Link-State vs Distance-Vector Algorithms

### 2.5.3 5.2 Section-Wide Summary

#### Key Terms and Definitions

- **Routing Algorithm:** The control-plane process for finding “good” (least-cost) paths through a network of routers.
- **Graph ( $G = (N, E)$ ):** Abstract model of a network where nodes  $N$  are routers and edges  $E$  are links with associated costs.
- **Least-Cost Path:** The path between two nodes with the minimum sum of edge costs.
- **Centralized Algorithm:** Requires complete, global network state (a full map) as input.
- **Decentralized Algorithm:** Nodes compute paths iteratively using only local information from neighbors.
- **Link-State (LS) Algorithm:** Canonical centralized approach (e.g., Dijkstra’s). Each node floods its local link information to all others.

Feature	Link-State (LS) Algorithm	Distance-Vector (DV) Algorithm
Information Shared	Each node broadcasts the cost of its directly attached links to all other nodes (flooding).	Each node sends its distance vector (its current cost estimates to all destinations) only to its immediate neighbors.
Knowledge	<b>Global:</b> Every node builds a complete map of the network topology.	<b>Local:</b> A node knows only its neighbors and the information received from them.
Message Complexity	High. Each link-state update is flooded across the entire network: typically $O(N E)$ messages.	Lower. Only neighbor-to-neighbor exchanges; typically $O(N^2)$ over time depending on convergence.
Speed of Convergence	Fast. Once link-state information propagates, each node runs Dijkstra in $O(N^2)$ or $O(N \log N)$ . No loops during convergence.	Slower. Can suffer from routing loops and the <i>count-to-infinity</i> problem.
Robustness	High. A node computes its own routes independently; a faulty node can only report incorrect local link costs.	Lower. Incorrect distance vectors can propagate from one node to another and corrupt the entire network state.
Used By	OSPF, IS-IS	RIP, BGP (path vector variant)

- **Dijkstra's Algorithm:** Greedy, iterative algorithm used in LS to compute shortest paths from a source to all other nodes.
- **Distance-Vector (DV) Algorithm:** Canonical decentralized approach (e.g., Bellman–Ford). Each node sends its distance vector to its neighbors.
- **Bellman–Ford Equation:**

$$d_x(y) = \min_{v \in N(x)} \{c(x, v) + d_v(y)\}$$

- **Count-to-Infinity:** A pathological DV behavior where a link cost increase causes slow propagation of “bad news,” leading to routing loops and repeatedly increasing distance estimates.
- **Poisoned Reverse:** A mitigation for 2-node loops in DV; a node advertises an infinite distance to a destination through the neighbor it routes through.

## Core Relationships

- Routing algorithms compute the paths that populate the forwarding tables in the data plane.

- Centralized algorithms correspond to LS (e.g., Dijkstra's).
- Decentralized algorithms correspond to DV (e.g., Bellman–Ford).
- The Bellman–Ford equation defines shortest paths recursively; DV implements this recursively and distributively.
- Count-to-Infinity is the major weakness of DV and results from slow propagation of link-cost increases.

### Key Insights and Takeaways

- All routing algorithms model the network as a graph and attempt to solve the least-cost path problem.
- **Fundamental tradeoff:**
  - LS (centralized): robust, fast convergence, but high message overhead (flooding).
  - DV (decentralized): low message overhead (neighbor-only), but slow convergence and vulnerable to loops.
- Pathological behaviors matter:
  - LS can oscillate when link costs depend on load.
  - DV can count to infinity when link costs increase.
- The Internet uses both approaches:
  - OSPF → LS
  - BGP → DV-like (path vector)
- No single algorithm is universally superior—both excel in different contexts.

## 3 5.3 Intra-AS Routing in the Internet: OSPF

### 3.1 Need for Routing Hierarchy: Scale and Autonomy

A "flat" network where every router runs same algorithm doesn't work for global internet for 2 reasons

1. Scale
  - (a) Storage and Computation
    - i. Internet has hundreds of millions of routers
    - ii. A link-state algorithm like Dijkstra's, which requires each router to store a map of the entire network and run an  $O(N^2)$  or  $O(N \log N)$  computation, is completely unfeasible at this scale.

- (b) COnmunication Overhead
  - i. a link-state algo requires link-state 'advertisements' (LSAs) to be flooded to every other router
  - ii. Communication overheard would be overwhelming.
  - iii. Distance vector would never converge
- 2. Administrative Autonomy
  - (a) Internet is not one network, it's composed of thousands of individual ISPs, universities, and corporate networks
  - (b) An ISP is its own administrative domain. It wants and needs to run its network as it sees fit. SHould be able to choose its own internal routing algo and manage its network for its own performance and economic goals.
  - (c) Furthermore, an ISP has no desire to "broadcast" details of its internal network topology and link costs to its competitors

### 3.2 Solution: Autonomous Systems (ASs)

- 1. Is a group of routers that are all under same admin control, typically a single ISP or a large organization
- 2. Each AS identified by a globally unique Autonomous System Number (ASN), which is assigned by ICANN regional registries
- 3. Two-level hierarchy
  - (a) Intra-AS Routing (Interior Gateway Protocol - IGP)
    - i. Routing algorithm that runs within a single AS
    - ii. All routes in AS are teammates and run samw intra-AS protocol
    - iii. Focus is on performance, finding best path within network
  - (b) Inter-AS Routing (Exterior Gateway Protocol - EGP)
    - i. routing algorithm that runs between ASs
    - ii. Protocol is used to route traffic across different ISPs
    - iii. Focus is on policy, enforcing economic and security rules (e.g. "I will not carry traffic from this competitor for free")

### 3.3 OSPF: Canonical Intra-AS Protocol

- 1. Open Shortest Path First
- 2. is a link-state protocol that uses Dijkstras algo to compute shortest paths
- 3. How it works

- (a) Construct map: each router in AS constructs a complete, identical topological map of entire AS
- (b) Flood Information: To build map, each router broadcasts its link-state information (i.e. its list of directly connected neighbors and costs of those links) to all other routers in AS.
- (c) Run Dijkstras: Once router has complete map, it runs Dijkstras alg using itself as root node. Finds least-cost path from itself to every other node (and subnet) in AS
- (d) Install Routes: router uses results of Dijkstras to build forwarding table

#### 4. Practical View: Link Costs

- (a) Cost of a link is not standard, its a configured parameter set by network admin
- (b) OSPF provides mechanism to find least-cost path, but policy of what "cost" means is up to admin

### 3.4 OSPF Link-State Broadcasts

Detail of OSPF is how it manages its link-state advertisements (LSAs)

- 1. Transport: messages are not sent over TCP or UDP, sent directly inside IP datagrams, using IP protocol number 89
- 2. Implication: Doesn't use a transport protocol, must implement its own reliability. Checks that links are operational (using HELLO messages sent to neighbors), and allows routers to get a neighbors link-state database
- 3. When to broadcast
  - (a) On Change: A router broadcasts its link-state information whenever one of its links changes state (e.g. goes down, comes up, or its cost is unchanged)
  - (b) Periodically: for robustness, router also re-broadcasts its link state at a regular interval even if there has been no change

### 3.5 Advanced Features in OSPFv2

is not just a simple implementation of Dijkstras. [RFC 2328] specifies several advances features that are critical in real-world networks

- 1. Security
  - (a) How do you stop a malicious router from joining AS and injecting false link-state advertisements?
  - (b) OSPF allows for authentication of OSPF messages
    - i. Simple Auth: clear-text password, which is weak

- ii. MD5 Auth
  - A. All routers configured with a shared secret key
  - B. Router computes an MD5 hash of OSPF packet (including key) and sends hash with packet
  - C. Receiver does same computation and verifies hash
  - D. Provides message integrity and authentication, and sequence numbers are used to prevent replay attacks
2. Multiple Same-Cost Paths
  - (a) What if Dijkstra's finds two paths to a destination with exact same cost?
  - (b) Instead of just picking one, OSPF allows router to install both paths in its forwarding table
  - (c) Traffic destined for that node can then be split across both paths
  - (d) Simple and effective form of load balancing
3. Support for Hierarchy (Areas)
  - (a) For a very large AS, flooding LSAs to every router is too much overhead
  - (b) Solution: AS can be partitioned into areas
  - (c) Intra-Area routing: routers within an area run their own OSPF algorithm. They flood LSAs only to other routers in same area. Each router in area has a complete map of only its area
  - (d) Backbone area (Area 0): One special area is designated as backbone. Primary role of backbone is to route traffic between different areas
  - (e) Area Border Routers (ABRs): Routers that sit at edge of an area. They're in backbone and in one or more other areas. Summarize routing information from their area and advertise it to the backbone, and vice-versa
4. Multicast support: OSPF extended by Multicast OSPF (MOSPF), which uses same link-state database and flooding mechanism to compute multicast paths

### 3.6 Section-Wide Summary

#### Key Terms and Definitions

- **Autonomous System (AS):** A group of routers under a single administrative control.
- **Intra-AS Routing (IGP):** Routing protocol used within an AS; focus is on performance.
- **Inter-AS Routing (EGP):** Routing protocol used between ASs; focus is on policy.
- **OSPF (Open Shortest Path First):** A widely used, open-standard, link-state, intra-AS routing protocol.

- **Link-State Advertisement (LSA):** A broadcast packet containing the state and cost of a router's links, used by OSPF.
- **Traffic Engineering:** Setting link costs to achieve network-wide performance goals (e.g., minimize congestion).
- **OSPF Area:** A sub-domain within a large AS used to create hierarchy and limit LSA flooding.
- **Backbone Area (Area 0):** The central area in a hierarchical OSPF network connecting all other areas.
- **Area Border Router (ABR):** A router that connects one or more areas to the backbone area.

### Core Relationships

- The Internet is divided into ASs to address issues of scale and administrative autonomy.
- This leads to a two-tier routing system: Intra-AS (e.g., OSPF) and Inter-AS (e.g., BGP).
- OSPF is a practical implementation of the Link-State algorithm (using Dijkstra's).
- Routers flood LSAs to build an identical network map, then compute shortest paths independently.
- Large ASs use OSPF hierarchy (Areas + Backbone) to limit LSA flooding and reduce computation.

### Key Insights / Takeaways

- Internet routing is fundamentally hierarchical: Intra-AS and Inter-AS levels enable global scalability.
- OSPF is the canonical Intra-AS protocol, relying on the Link-State approach.
- OSPF offers robustness: authentication, multi-path load balancing, and hierarchical scalability.
- OSPF link costs are policy tools—not purely physical measurements—and can be tuned for Traffic Engineering.

## 4 5.4 Routing Among the ISPs: BGP

With OSPF, a link-state protocol, is designed for scalability and performance within a single administrative domain. All routes within AS are teammates, running the same protocol with the shared goal of finding the best-performing, least-cost paths.

### 1. BGP

- (a) Fundamentally different from OSPF
- (b) OSPF goal is performance, BGP's goal is policy
- (c) Is a decentralized, async protocol that is similar in many ways to Distance-Vector algo

### 4.1 5.4.1 Role of BGP

- 1. Why do we need a separate protocol for inter-AS routing?
  - (a) Reasons are scale and autonomy
  - (b) OSPF's 'everyone knows the full map' approach won't scale to hundreds of millions of routers on global internet
  - (c) Furthermore, an AS doesn't want to share its internal topology with competitors
- 2. Fundamentals of BGP
  - (a) Within an AS, a router uses its intra-AS protocol (OSPF) to find best path to any other router inside of AS1 (let's say router 1a in AS1), but what about a destination outside its AS, say a prefix x located in AS3?
  - (b) Router 1a's OSPF built forwarding table is useless for this. It needs mechanisms to:
    - i. Discover that prefix x even exists
    - ii. Learn a path to prefix x
    - iii. Choose best path from potentially many options, based on its AS's policies
- 3. Routing to Prefixes
  - (a) Key distinction with BGP is that it doesn't route individual host IP addresses
  - (b) Routes to prefixes, which are CIDRized blocks of addresses (e.g. 138.16.18/22) that represent a subnet or a collection of subnets
  - (c) Every router forwarding table is thus a combination of routes learned from two different sources
    - i. Intra-AS routes (from OSPF): typically for destinations inside the AS
    - ii. Inter-AS routes (from BGP): for destinations outside AS
  - (d) BGP provides each router with means to:

- i. Obtain prefix reachability information. BGP allows an AS to advertise its prefixes to rest of internet and ensures this advertisement propagates to all other ASs
- ii. Determine best route: a router might learn multiple paths to same prefix. For example, AS1 might learn it can reach x via AS2, or via direct link to AS3

## 4.2 5.4.2 Advertising BGP Route Information

### 1. Concept Overview: iBGP and eBGP

How does a prefix advertisement (e.g., “AS3 has x”) get from a router deep inside AS3 to a router deep inside AS1?

The process involves two “flavors” of BGP:

- **External BGP (eBGP):** Used to communicate between routers in different ASs.
- **Internal BGP (iBGP):** Used to communicate between routers within the same AS.

Before we see how they work, we must define two types of routers:

- **Gateway Router:** A router at the “edge” of an AS that has a direct physical link to a router in another AS (e.g., router 1c in Figure 5.8).
- **Internal Router:** A router that is connected only to other hosts and routers within its own AS (e.g., router 1b in Figure 5.8).

### 2. Technical Mechanism: BGP Connections

Unlike OSPF, BGP does not send its messages raw over IP. BGP runs over TCP on port 179.

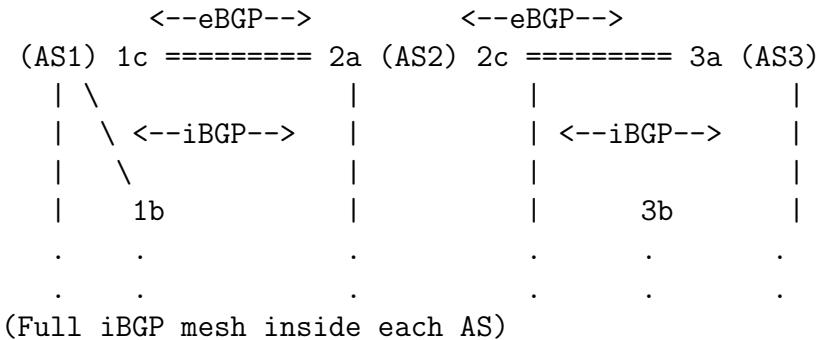
A BGP connection (or BGP session) is the semi-permanent TCP connection between two routers running BGP.

An eBGP connection is a BGP session that spans two ASs (e.g., between router 1c in AS1 and router 2a in AS2).

An iBGP connection is a BGP session between routers in the same AS (e.g., between 1c and 1b in AS1).

This is a critical point: iBGP connections are logical connections, not necessarily physical. As shown in Figure 5.9, the iBGP session between 1c and 1b runs over the AS1 network. The TCP packets that carry the BGP message from 1c to 1b are themselves routed by AS1’s intra-AS protocol (OSPF).

**Textual Representation of Figure 5.9: BGP Connections**



### 3. Practical View: The Advertisement Propagation Chain

Let's trace the advertisement for prefix x from AS3 to all routers in AS1 (Figure 5.8/5.9):

- (a) **Start:** Router 3a (gateway) in AS3 advertises x to router 2c (gateway) in AS2 via an eBGP connection. The message is effectively: “I can reach x”.
  - (b) **Internal Propagation (AS2):** Router 2c receives this advertisement. It must now inform all other routers in its own AS (AS2) about this new prefix. It sends an iBGP message to all its iBGP peers (e.g., 2a, 2b, 2d) saying “I can reach x via 3a”.
  - (c) **Cross-AS Propagation:** Router 2a (gateway) now knows it can reach x. It advertises this fact to router 1c (gateway) in AS1 via their eBGP connection. The message is now: “I can reach x, and the path is (AS2, AS3)”.
  - (d) **Internal Propagation (AS1):** Router 1c receives this advertisement. It uses iBGP to inform all its internal peers (1a, 1b, 1d) of this new route: “I can reach x via the path (AS2, AS3)”.

**End State:** Now, every router in AS1 and AS2 knows about prefix x and also knows the path of ASes required to get there.

### 4.3 5.4.3 Determining the Best Routes

## 1. Concept Overview: BGP Routes and Attributes

In the previous example, AS1 learned one path to x. But in the real Internet (as in Figure 5.10), AS1 might learn multiple paths to x. For example, if AS1 also has a direct link to AS3, it might learn two paths:

Path 1: (AS2, AS3, x)

Path 2: (AS3, x)

How does a router in AS1 choose? It uses BGP's route-selection algorithm.

A route in BGP is more than just a prefix; it's the prefix plus its associated attributes. Two of the most important attributes are AS-PATH and NEXT-HOP.

## 2. Technical Mechanism: AS-PATH and NEXT-HOP Attributes

### AS-PATH:

- **What it is:** This attribute contains the list of ASes that the advertisement has traversed.
- **How it's built:** When a prefix is advertised outside an AS, the AS prepends its own ASN to the list.
  - 3a advertises (AS3, x) to 2c.
  - 2c propagates this internally.
  - 2a advertises to 1c, prepending its own AS: (AS2, AS3, x).
- **Function 1 (Loop Detection):** This is BGP's primary loop-detection mechanism. If a router receives an advertisement where its own AS is already in the AS-PATH, it rejects the route.
- **Function 2 (Metric):** The length of the AS-PATH (number of ASes) is used as a routing metric. Shorter paths are preferred.

### NEXT-HOP:

- **What it is:** This is the IP address of the router interface that begins the AS-PATH.
- **Crucial Role:** This attribute is the critical link that connects the inter-AS (BGP) routing logic to the intra-AS (OSPF) forwarding logic.

Example (Figure 5.10):

Router 1c learns the route (AS2, AS3, x) from router 2a. The NEXT-HOP for this route is the IP address of 2a's interface.

Router 1d learns the route (AS3, x) from 3d. The NEXT-HOP for this route is the IP address of 3d's interface.

When 1c and 1d advertise these routes internally via iBGP to router 1b, they include the original NEXT-HOP attribute.

Router 1b's (internal) view:

Route 1 for x: AS-PATH = (AS2, AS3), NEXT-HOP = <IP of 2a>

Route 2 for x: AS-PATH = (AS3), NEXT-HOP = <IP of 3d>

## 3. Technical Mechanism: Hot Potato Routing

Given these two routes, how does 1b choose? The simplest policy is Hot Potato Routing.

**Concept:** This is a “selfish” algorithm. The AS wants to get the packet “off its network” (like a hot potato) as quickly and cheaply as possible. It doesn't care about the cost outside its own network.

Mechanism (Figure 5.11):

- Router 1b looks at the two routes it knows for x.
- It examines their NEXT-HOP attributes: <IP of 2a> and <IP of 3d>.
- 1b then consults its intra-AS (OSPF) routing table to find the internal cost to reach <IP of 2a> and the internal cost to reach <IP of 3d>.
- It selects the route whose NEXT-HOP is “closer” based on its internal OSPF costs.

If 2a is 2 hops away and 3d is 3 hops away, 1b chooses Route 1. It then looks at its OSPF forwarding table to find the local interface I that is on the shortest path to 2a, and it installs (x, I) in its forwarding table.

**Key Insight:** Hot Potato Routing highlights the symbiosis of BGP and OSPF. BGP provides the external path options, and OSPF provides the internal costs to reach those external paths.

#### 4. Technical Mechanism: The Full BGP Route-Selection Algorithm

Hot Potato Routing is just one step in the full algorithm. BGP routers are driven by policy. The full algorithm is a sequential list of rules. A router takes all known routes for a prefix and applies these rules in order until only one route is left.

- **Local Preference (POLICY):** Each route is assigned a local\_pref value by the router that learns it. This is a policy decision. The router must select the route with the highest local preference value, regardless of any other attribute.
- **Shortest AS-PATH (PERFORMANCE):** From the remaining routes (which are tied for the highest local\_pref), select the route with the shortest AS-PATH length.
- **Hot Potato Routing (COST):** From the remaining routes (which are tied on local\_pref and AS-PATH length), select the route with the closest NEXT-HOP router (based on intra-AS (OSPF) costs).
- **BGP Identifiers:** If still tied, use other BGP attributes (like router IDs) as a final tie-breaker.

Example Revisited: In Figure 5.10, 1b learns (AS2, AS3, x) and (AS3, x).

Assume local\_pref is the same for both (Rule 1 tie).

Rule 2 (Shortest AS-PATH) is applied. (AS3, x) has length 1. (AS2, AS3, x) has length 2.

Router 1b selects the route (AS3, x).

Notice that Hot Potato Routing (Rule 3) is not even used in this case, because the AS-PATH lengths were different. This shows BGP is not a purely selfish protocol; it does prefer shorter AS-level paths, unless policy (Rule 1) dictates otherwise.

#### 4.4 5.4.4 IP-Anycast

BGP's "best path" selection logic enables a powerful service called IP-Anycast, which is heavily used by the DNS system.

**Concept:** Replicate a service (like a DNS root server) in many different geographical locations, but assign the exact same IP address to all of them.

**Mechanism (Figure 5.12):**

- The CDN server in AS1 advertises 212.21.21.21 via BGP.
- The CDN server in AS4 also advertises the same prefix 212.21.21.21 via BGP.
- A router in AS3 will receive two different routes for 212.21.21.21: one via AS1 and one via AS4.
- The router in AS3 simply runs its standard BGP route-selection algorithm. It will pick the "best" path (e.g., the one with the shortest AS-PATH) and install that route in its forwarding table.

**Result:** When a client in AS3 sends a packet to 212.21.21.21, the BGP-built forwarding tables will automatically route the packet to the "closest" server instance. This provides incredible robustness and low latency for critical services like DNS.

#### 4.5 5.4.5 Routing Policy

As we've stressed, policy is the primary driver of BGP. The examples in Figure 5.13 illustrate the two most common policy relationships: provider-customer and peer-peer.

**Policy Example 1: Stub Networks (Customer/Provider)**

- **Scenario:** X is a customer of B and C. W and Y are customers of A and C, respectively.
- **Goal:** X is a "stub" network. It should originate and terminate traffic, but it should never be a transit AS. That is, B should not send traffic destined for C through X.
- **Implementation (Policy):** X does not advertise routes learned from one provider (C) to its other provider (B).
- **Result:** Since B never learns a path to Y or C via X, it will never route traffic through X to get there. Policy is enforced by selectively choosing what routes to advertise.

**Policy Example 2: Peering (No "Free Rides")**

- **Scenario:** B and C are peers. A is a customer of B.
- **Goal:** B and C agree to exchange traffic between their respective customers. B does not want to carry traffic from C to the rest of the Internet (e.g., to A) for free.
- **Implementation (Policy):**
  - B learns a route to W via its customer A (path A-W).

- B will advertise this route (B-A-W) to its customer X (because X pays B for Internet transit).
  - B will *not* advertise this route (B-A-W) to its peer C.
- **Result:** C never learns a path to W via B. C cannot use B as a “free ride” to get to A’s customers.

**Rule of Thumb:** An ISP only carries traffic if it originates from or is destined to one of its own customers.

**Sidebar: Why Different Inter-AS and Intra-AS Protocols?** **Policy:** Inter-AS (BGP) is all about policy. It allows an AS to control what traffic it carries and what routes it advertises based on business relationships (customer, peer). Intra-AS (OSPF) has no concept of this; its only goal is to find the best technical path.

**Scale:** OSPF’s “everyone-knows-the-map” (link-state flooding) approach does not scale to the hundreds of thousands of networks on the Internet. BGP (a DV-like protocol) scales by only advertising paths of ASes, not individual router links.

**Performance:** Intra-AS (OSPF) can be finely tuned for performance, using link costs to optimize latency or bandwidth. Inter-AS (BGP) often sacrifices performance for policy. A longer, more expensive path that satisfies a business agreement will be chosen over a shorter one that doesn’t.

#### 4.6 5.4.6 Putting the Pieces Together: Obtaining Internet Presence

This section provides a perfect case study integrating everything from Chapter 2, 4, and 5.

**Goal:** Your new company, Xanadu Inc., wants to connect its network (Web, mail, DNS servers) to the global Internet.

##### The Protocol Checklist:

- **Get Physical:** Contract with a local ISP. You get a physical link and a block of IP addresses (a CIDR prefix, e.g., /24) from your ISP.
- **Configure Local Network (IP/DHCP):** You assign IP addresses from your block to your gateway router and your servers. You might use DHCP (Chapter 4) to assign addresses to your internal employee hosts.
- **Get a Name (DNS):** You register your domain name (e.g., xanadu.com) with a DNS registrar.
- **Register Your DNS Server (DNS):** You must tell the registrar the name and IP address of your company’s authoritative DNS server. The registrar then inserts an NS record (Name Server record) into the .com TLD servers, “delegating” xanadu.com to your server’s IP address.
- **Populate Your DNS Server (DNS):** In your own DNS server, you create the resource records for your company:

- An A record mapping `www.xanadu.com` to your Web server’s IP address.
- An MX record mapping `xanadu.com` to your mail server’s name.
- **Get Global Reachability (BGP):** This is the final, essential step. Your ISP’s gateway router must advertise your prefix (your /24 block) to its upstream providers using eBGP. This advertisement then propagates via BGP across the entire Internet.

### The Result (A Day in the Life):

- When a user types `www.xanadu.com`:
  - Their DNS query (Chapter 2) goes to a root server, then a .com TLD server.
  - The TLD server sees the NS record (Step 4) and directs the query to your DNS server.
  - Your DNS server replies with the A record (Step 5), giving the user your Web server’s IP.
- When the user’s browser sends a TCP SYN (Chapter 3) to that IP:
  - Every router on the Internet uses its forwarding table—built by BGP (Step 6)—to find the AS-PATH toward your ISP, and ultimately, to your network.
  - The packet arrives at your gateway router.
  - Your gateway router uses its internal (e.g., OSPF) forwarding table to send the packet to the final destination: your Web server.

## 4.7 5.4 Section-Wide Summary

### Key Terms and Definitions:

- **BGP (Border Gateway Protocol):** The de facto inter-AS (inter-domain) routing protocol for the global Internet.
- **Inter-AS vs. Intra-AS:** Inter-AS is between autonomous systems (policy-driven); Intra-AS is within an autonomous system (performance-driven).
- **Gateway Router:** An AS-edge router that connects to a router in another AS.
- **Internal Router:** A router that only connects to devices within its own AS.
- **eBGP:** The BGP session between two different ASs. Used to advertise external routes.
- **iBGP:** The BGP session within an AS. Used to propagate external routes to all internal routers.
- **BGP Connection:** A logical session running on TCP port 179.
- **Route:** In BGP, a route is a prefix + attributes.

- **AS-PATH Attribute:** The list of ASes an advertisement has traversed. Used for loop detection and as a metric (shorter is better).
- **NEXT-HOP Attribute:** The IP address of the next router on the external path. This is the crucial link between BGP (inter-AS) and OSPF (intra-AS).
- **Hot Potato Routing:** A “selfish” routing policy where an AS forwards a packet to the internally closest exit point (NEXT-HOP), regardless of the external path cost.
- **Route Selection Algorithm:** The 4-step process BGP uses to select a single “best” route: 1. Highest Local Preference (Policy), 2. Shortest AS-PATH, 3. Closest NEXT-HOP (Hot Potato), 4. Tie-breakers.
- **IP-Anycast:** A service (used by DNS) where a single IP address is assigned to multiple, geographically dispersed servers, and BGP’s route selection automatically directs clients to the “closest” one.

### Core Relationships:

- BGP is the “glue” of the Internet, connecting all the ASes.
- BGP uses eBGP to talk to other ASes and iBGP to inform its own internal routers.
- The BGP Route Selection Algorithm shows the priority of Policy (Local Preference) over Performance (AS-PATH, Hot Potato).
- A router’s final forwarding table is built using both BGP (to find the exit gateway/NEXT-HOP for external prefixes) and OSPF (to find the internal path to that NEXT-HOP).

### Key Insights / Takeaways:

- Inter-domain routing (BGP) is not about finding the best path; it’s about finding a policy-compliant path that works. Business relationships (customer/provider, peer/peer) dictate routing, not just link speeds.
- BGP is a Path Vector protocol, a “DV-like” protocol. Instead of just advertising a cost (distance), it advertises the entire path (the AS-PATH attribute). This is what allows it to implement policy and detect loops robustly.
- A router in an AS makes its “best” choice based on the information it has. Hot Potato Routing is the classic example of this “selfish” behavior: minimize my cost, and let the next AS worry about its cost.
- The Internet is a “network of networks.” To “join” the Internet, you need more than a physical link and an IP address; your AS must participate in BGP so the rest of the world can learn a route to your prefixes.

## 5 5.5 The SDN Control Plane

Logically centralized control, which is foundation of Software-Defined Networking (SDN)

1. Control plan is physically removed from routers and implemented in a separate, remote controller
2. Unbundling of network has implications for how networks are designed, managed, and programmed.

## 6 The Four Key Characteristics of an SDN Architecture

SDN is not a single protocol but an architectural approach defined by the following

### 6.1 Flow-based Forwarding

This is the data plane mechanism that SDN controls. We studied this in Chapter 4 as “generalized forwarding.”

Unlike traditional routing, which forwards packets based only on the destination IP address, SDN-controlled switches (which we call “packet switches”) use a flow table.

This flow table implements the “match-plus-action” paradigm. A “match” can be made on numerous header fields across multiple layers (e.g., L2 MAC address, L3 IP address, L4 TCP/UDP port).

The “action” is also general: it can be to forward, drop, modify, or send the packet to the controller. It is the SDN control plane’s job to compute, install, and manage these flow table entries.

### 6.2 Separation of Data Plane and Control Plane

This is the core philosophy of SDN.

The data plane consists of the network’s switches. These are now “dumb” (but fast) hardware devices whose only job is to execute the match-plus-action rules in their flow tables.

The control plane consists of the remote servers and software (the “brain”) that determine and manage these flow tables.

### 6.3 Network Control Functions are External to Data-Plane Switches

This highlights that the control plane is physically separate from the data-plane switches.

As shown in Figure 5.14, the control plane itself is composed of two parts:

- **The SDN Controller (or Network Operating System):** This is the core platform. It maintains a centralized view of the network’s state (links, switches, hosts) and provides an API for applications to interact with it.

- **Network-Control Applications:** These are the “apps” that run on the controller. They contain the actual logic for network functions like routing, load balancing, or firewalling. They use the controller’s API to specify the desired behavior.

The controller is “logically centralized,” meaning it presents a single, unified view of the network. In practice, it’s a physically distributed system (run on multiple servers) for fault tolerance and scalability.

## 6.4 A Programmable Network

This is the payoff of the first three characteristics. The network becomes programmable via the network-control applications.

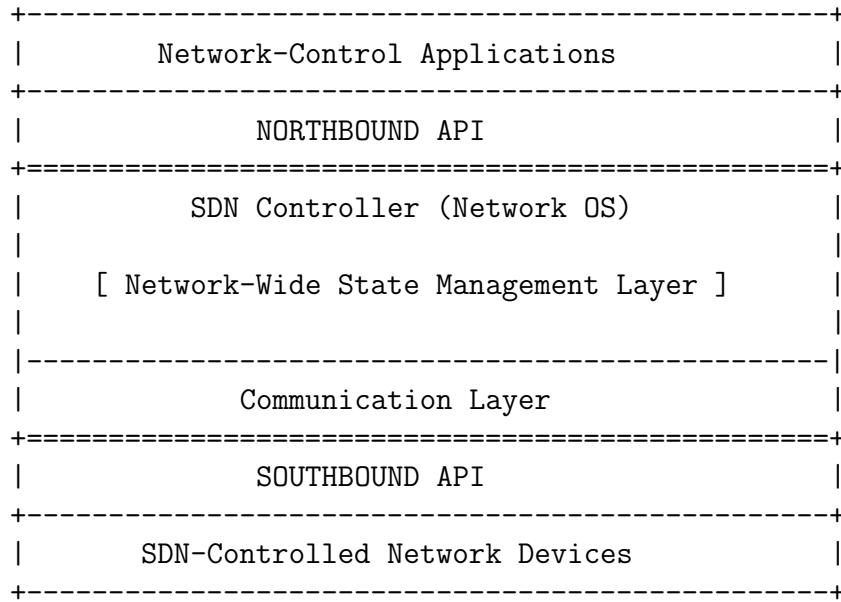
Analogy (The PC vs. Mainframe): The textbook correctly likens this to the “unbundling” of the PC from the monolithic mainframe.

Old Model (Mainframe/Traditional Router): A single vendor sold you a “vertically integrated” box containing the hardware, the operating system, and the applications.

New Model (PC/SDN): The architecture is “unbundled” and open. You can get data-plane switch hardware from one vendor, an SDN controller OS from another, and network-control apps from a third vendor, or write them yourself.

### 7 5.5.1 The SDN Control Plane: Controller and Applications

#### Textual Representation of Figure 5.15: SDN Controller Architecture



## 7.1 The Communication Layer (Southbound Interface)

Purpose: This layer provides the communication channel between the SDN controller and the network's data-plane switches.

Mechanism: Requires a standardized protocol for two-way communication.

Example Protocol: The most prominent southbound protocol is OpenFlow.

## 7.2 The Network-Wide State-Management Layer

Purpose: This is the controller's "single source of truth." It maintains a real-time, comprehensive view of the entire network's state.

Mechanism: A distributed database storing host, link, switch state, flow tables, and a topology graph.

## 7.3 The Interface to Applications (Northbound Interface)

Purpose: This is the API that applications use to interact with the controller.

Mechanism: Applications can read state, write state, and register for notifications.

Example APIs: Many controllers expose a RESTful API.

## 7.4 Practical Implementation: ODL and ONOS

### The OpenDaylight (ODL) Controller

Key Component: The Service Abstraction Layer (SAL) is the central "message bus."

Southbound: Supports OpenFlow, SNMP, NETCONF.

Northbound: AD-SAL (REST) and MD-SAL (YANG-based).

### The ONOS Controller

Key Component: The ONOS Distributed Core is a distributed system.

Northbound (Intent Framework): Applications specify high-level intents, and ONOS converts them into flow rules.

## 8 5.5.2 OpenFlow Protocol

### 8.1 Concept Overview

OpenFlow defines the communication between an SDN controller and switch over TCP (port 6653).

### 8.2 Technical Mechanism: Key Message Types

**Controller-to-Switch:** Configuration, Modify-State, Read-State, Send-Packet.

**Switch-to-Controller:** Flow-Removed, Port-status, Packet-in.

### 8.3 Practical View: Google's B4 Network

A global SDN WAN connecting Google data centers. Uses OpenFlow switches, distributed controllers, and a Traffic Engineering application.

## 9 5.5.3 Data and Control Plane Interaction: An Example

Scenario: A link between s1 and s2 fails.

Sequence of events: Data-plane failure → Port-status message → state update → application notified → Dijkstra recomputation → new flow rules installed → Modify-State messages sent → switches update tables.

## 10 5.5.4 SDN: Past and Future

### 10.1 Past: The Origins of SDN

Earlier proposals separated planes; Ethane (2007) led directly to OpenFlow.

### 10.2 Future: Network Functions Virtualization (NFV)

NFV replaces hardware middleboxes with software on commodity hardware.

### 10.3 Future: Inter-AS SDN

Research attempts to extend SDN to multi-ISP environments.

## 11 5.5 Section-Wide Summary

### 11.1 Key Terms and Definitions

SDN, Flow-Based Forwarding, SDN Controller, Network-Control Application, Southbound API, Northbound API, OpenFlow, NFV.

### 11.2 Core Relationships

SDN separates planes; generalized forwarding is the data-plane model; OpenFlow is the southbound protocol; applications use the northbound API; NFV applies unbundling to middleboxes.

### 11.3 Key Insights / Takeaways

SDN unbundles the router; controller is logically centralized but physically distributed; control plane is layered; intelligence moves to applications; Packet-in and Modify-State form the core reactive loop.

1. X

(a) Y