

Principles of Reliable Data Transfer

Study-Ready Notes

Compiled by Andrew Photinakis

October 17, 2025

Contents

1	Introduction to Reliable Data Transfer	2
2	Reliable Data Transfer Service Abstraction	2
2.1	Service Model	2
3	RDT Protocol Interfaces and Mechanisms	2
3.1	Protocol Interface Functions	2
3.2	Key Protocol Mechanisms	3
4	Finite State Machine Specification	3
4.1	FSM Concepts	3
5	RDT1.0: Reliable Transfer over Reliable Channel	3
5.1	Assumptions	3
5.2	FSM Specifications	4
6	RDT2.0: Channel with Bit Errors	4
6.1	New Challenges	4
6.2	Error Recovery Mechanisms	4
7	RDT2.1: Handling Garbled ACK/NAKs	4
7.1	The Fatal Flaw in RDT2.0	4
7.2	Solution: Sequence Numbers	5
8	RDT2.2: NAK-Free Protocol	5
8.1	Approach	5
9	RDT3.0: Channels with Errors and Loss	6
9.1	New Channel Assumptions	6
9.2	Solution: Countdown Timer	6

10 RDT3.0 Operation Scenarios	6
10.1 Four Key Scenarios	6
11 Performance Analysis of RDT3.0	7
11.1 Sender Utilization	7
11.2 Utilization Formula	7
11.3 Example Calculation	7
12 Pipelined Protocols	8
12.1 Limitations of Stop-and-Wait	8
12.2 Pipelining Solution	8
12.3 Utilization with Pipelining	8
13 Go-Back-N Protocol	8
13.1 Sender Operation	8
13.2 Sender Window Structure	9
14 Selective Repeat Protocol	9
14.1 Key Features	9
14.2 Window Size Constraint	9
15 Study Aids and Exam Preparation	10
15.1 Key Concepts to Master	10
15.2 Practice Questions	10
16 Summary	11

1 Introduction to Reliable Data Transfer

- Core component of transport-layer services
- Essential for connection-oriented protocols like TCP
- Builds upon multiplexing/demultiplexing concepts
- Addresses challenges of unreliable network channels

[Summary: Reliable data transfer protocols ensure data delivery despite unreliable network conditions, using mechanisms like acknowledgments, sequence numbers, and retransmissions.]

2 Reliable Data Transfer Service Abstraction

2.1 Service Model

- **Abstraction:** Perfectly reliable channel between sending and receiving processes
- **Reality:** Underlying channel is unreliable (loses, corrupts, reorders data)
- **Implementation:** Transport layer protocols handle reliability

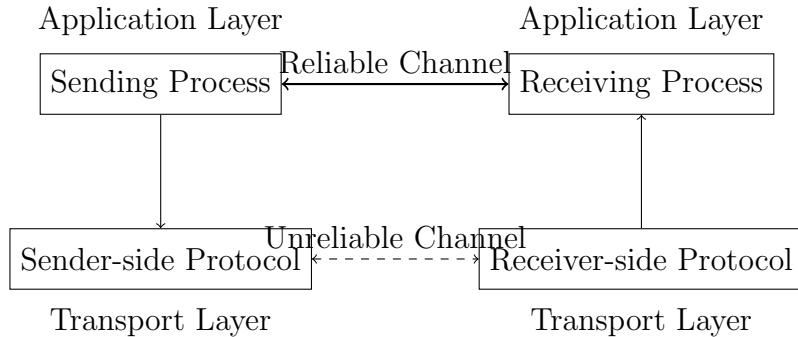


Figure 1: Reliable data transfer: abstraction vs implementation

3 RDT Protocol Interfaces and Mechanisms

3.1 Protocol Interface Functions

- **rdt_send():** Called from above to deliver data to receiver upper layer
- **udt_send():** Called by RDT to transfer packet over unreliable channel
- **rdt_rcv():** Called when packet arrives on receiver side
- **deliver_data():** Called by RDT to deliver data to upper layer

3.2 Key Protocol Mechanisms

- Error detection (checksums)
- Acknowledgments (ACKs) and Negative Acknowledgments (NAKs)
- Retransmission
- Sequence numbers for duplicate detection
- Timers for loss detection

[Mnemonic: "ACE RT" - Acknowledgments, Checksums, Error detection, Retransmission, Timers - key RDT mechanisms.]

4 Finite State Machine Specification

4.1 FSM Concepts

- Used to specify sender and receiver behavior
- **State:** Current condition of protocol entity
- **Event:** Trigger that may cause state transition
- **Actions:** Operations performed during transition

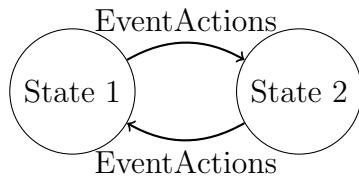


Figure 2: Finite State Machine basic structure

5 RDT1.0: Reliable Transfer over Reliable Channel

5.1 Assumptions

- Underlying channel is perfectly reliable
- No bit errors
- No loss of packets

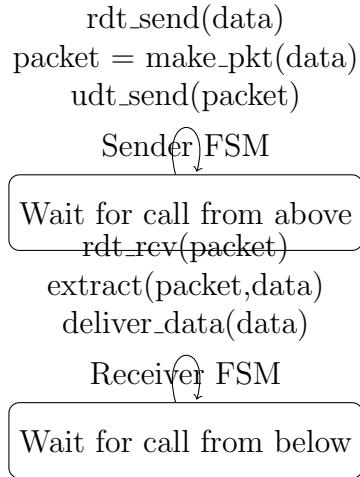


Figure 3: RDT1.0: Simple FSMs for perfectly reliable channel

5.2 FSM Specifications

[Summary: RDT1.0 assumes perfect channel conditions - sender simply sends packets, receiver simply receives and delivers them without error handling.]

6 RDT2.0: Channel with Bit Errors

6.1 New Challenges

- Underlying channel may flip bits in packets
- Use checksums to detect bit errors
- Recovery mechanisms needed

6.2 Error Recovery Mechanisms

- **Acknowledgments (ACKs)**: Receiver tells sender packet received OK
- **Negative Acknowledgments (NAKs)**: Receiver tells sender packet had errors
- **Stop-and-wait**: Sender sends one packet, waits for receiver response

7 RDT2.1: Handling Garbled ACK/NAKs

7.1 The Fatal Flaw in RDT2.0

- What if ACK/NAK becomes corrupted?
- Sender doesn't know what happened at receiver

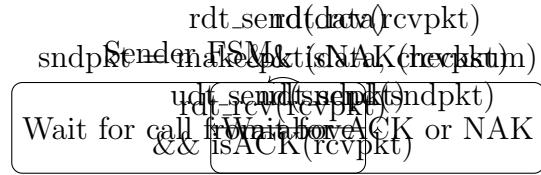


Figure 4: RDT2.0 sender FSM handling bit errors

- Can't just retransmit: possible duplicate

7.2 Solution: Sequence Numbers

- Add sequence number to each packet
- Two sequence numbers (0,1) suffice for stop-and-wait
- Receiver discards duplicate packets

[Concept Map: RDT Evolution → RDT1.0 (perfect channel) → RDT2.0 (bit errors, ACK/NAK) → RDT2.1 (garbled ACKs, sequence numbers) → Each version adds mechanisms to handle more realistic channel impairments.]

8 RDT2.2: NAK-Free Protocol

8.1 Approach

- Same functionality as RDT2.1 but using ACKs only
- Instead of NAK, receiver sends ACK for last packet received OK
- Receiver explicitly includes sequence number of packet being ACKed
- Duplicate ACK at sender triggers retransmission (same as NAK)

RDT2.1 (with NAKs)	RDT2.2 (NAK-free)
Uses ACK and NAK messages	Uses only ACK messages
Receiver sends NAK for corrupted packets	Receiver sends ACK for last good packet
Simpler logic	Reduced message types
Two response types	Unified response mechanism
Used in: Educational examples	Used in: TCP (real-world)

Table 1: Comparison of RDT2.1 vs RDT2.2

9 RDT3.0: Channels with Errors and Loss

9.1 New Channel Assumptions

- Underlying channel can lose packets (data and ACKs)
- Need mechanism to detect and recover from lost packets

9.2 Solution: Countdown Timer

- Sender waits "reasonable" amount of time for ACK
- Retransmits if no ACK received in this time
- If packet just delayed (not lost): retransmission creates duplicate
- Sequence numbers handle duplicates

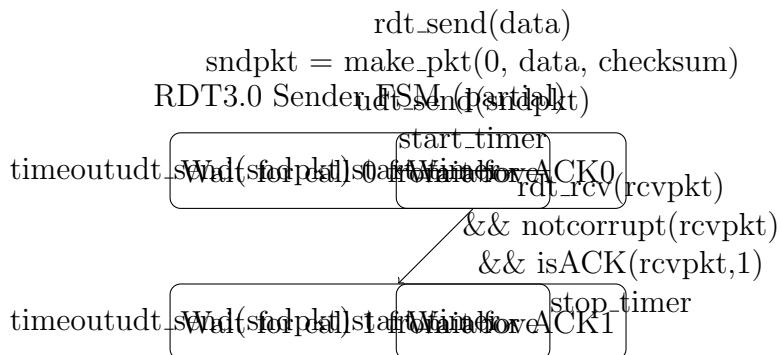


Figure 5: RDT3.0 sender with timer for loss recovery

10 RDT3.0 Operation Scenarios

10.1 Four Key Scenarios

1. **No loss:** Normal operation with timely ACKs
2. **Packet loss:** Timeout triggers retransmission
3. **ACK loss:** Duplicate ACK triggers retransmission
4. **Premature timeout:** Delayed ACK causes unnecessary retransmission

[Summary: RDT3.0 handles both bit errors and packet loss using checksums, sequence numbers, acknowledgments, and countdown timers, providing complete reliability over unreliable channels.]

Scenario	Problem	Solution
No loss	None	Normal operation
Packet loss	Data packet lost in transit	Timeout and retransmission
ACK loss	Acknowledgment lost	Timeout and retransmission
Premature timeout	ACK delayed, not lost	Sequence numbers handle duplicate

Table 2: RDT3.0 operation scenarios and solutions

11 Performance Analysis of RDT3.0

11.1 Sender Utilization

- U_{sender} : Fraction of time sender is busy sending
- For stop-and-wait protocol: very low utilization

11.2 Utilization Formula

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R}$$

Where:

- L = packet length (bits)
- R = transmission rate (bps)
- RTT = round-trip time

11.3 Example Calculation

- 1 Gbps link, 15 ms propagation delay, 8000 bit packet
- Transmission time: $D_{\text{trans}} = \frac{L}{R} = \frac{8000}{10^9} = 8$ microseconds
- $RTT = 2 \times 15$ ms = 30 ms
- $U_{\text{sender}} = \frac{0.008}{30.008} = 0.00027$

[Mnemonic: "Low Utilization Long Wait" - Stop-and-wait protocols have low utilization due to waiting for ACKs.]

12 Pipelined Protocols

12.1 Limitations of Stop-and-Wait

- Very low utilization (0.027% in example)
- Protocol limits performance of underlying infrastructure
- Most of the time, sender is idle waiting for ACK

12.2 Pipelining Solution

- Allow multiple "in-flight" packets
- Increase range of sequence numbers
- Require buffering at sender and/or receiver
- Dramatically increases utilization

12.3 Utilization with Pipelining

$$U_{\text{sender}} = \frac{N \times L/R}{RTT + L/R}$$

Where N is the number of pipelined packets.

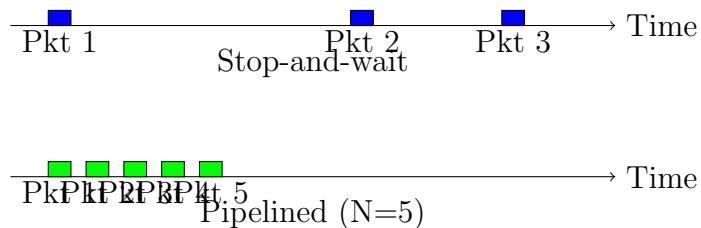


Figure 6: Comparison of stop-and-wait vs pipelined transmission

13 Go-Back-N Protocol

13.1 Sender Operation

- Window of up to N consecutive transmitted but unACKed packets
- k -bit sequence number in packet header
- **Cumulative ACK:** $\text{ACK}(n)$ acknowledges all packets up to and including sequence number n

- Single timer for oldest in-flight packet
- Timeout(n): retransmit packet n and all higher sequence number packets in window

13.2 Sender Window Structure

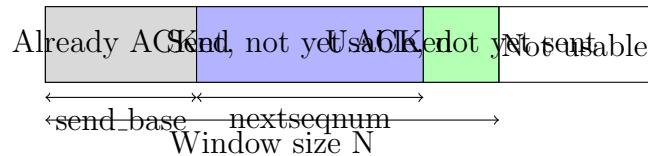


Figure 7: Go-Back-N sender window structure

14 Selective Repeat Protocol

14.1 Key Features

- Receiver individually acknowledges all correctly received packets
- Buffers out-of-order packets for eventual in-order delivery
- Sender maintains timer for each unACKed packet
- More efficient but more complex than Go-Back-N

14.2 Window Size Constraint

- Critical relationship between sequence number space and window size
- For sequence numbers 0 to $2^k - 1$, window size must satisfy:

$$N \leq 2^{k-1}$$

- Prevents ambiguity between old and new packets

[Concept Map: Pipelined Protocols → Go-Back-N (cumulative ACKs, simple) vs Selective Repeat (individual ACKs, efficient) → Trade-off between complexity and performance → Window size constraints critical for correctness.]

Go-Back-N	Selective Repeat
Cumulative acknowledgments	Individual acknowledgments
Discards out-of-order packets	Buffers out-of-order packets
Single timer for window	Timer per packet
Simpler implementation	More complex implementation
Less efficient with loss	More efficient with loss
Window: $N \leq 2^k - 1$	Window: $N \leq 2^{k-1}$

Table 3: Comparison of Go-Back-N vs Selective Repeat

15 Study Aids and Exam Preparation

15.1 Key Concepts to Master

- Understand the evolution of RDT protocols and what problem each version solves
- Be able to draw and interpret FSM diagrams for each protocol version
- Calculate sender utilization for stop-and-wait and pipelined protocols
- Explain the window size constraints for Go-Back-N and Selective Repeat
- Compare and contrast different reliability mechanisms

15.2 Practice Questions

1. **Trace the evolution** of RDT protocols from 1.0 to 3.0. What specific problem does each new version address, and what mechanism does it introduce?
2. Calculate the **sender utilization** for a stop-and-wait protocol with the following parameters: packet size = 1500 bytes, transmission rate = 100 Mbps, RTT = 50 ms. How does this change if we use pipelining with a window size of 10?
3. Explain why **Selective Repeat** requires that the window size be at most half the sequence number space, while **Go-Back-N** can use almost the entire sequence number space.
4. Draw the **FSM for RDT3.0 sender** and explain how it handles packet loss, ACK loss, and premature timeouts.
5. Compare the **performance characteristics** of Go-Back-N vs Selective Repeat in a network with high packet loss rates. Which would you choose and why?

[Mnemonic: "RDT 1-2-3" - RDT1.0 (perfect), RDT2.0 (errors), RDT3.0 (loss) - each adds one more real-world problem to solve.]

16 Summary

- Reliable data transfer protocols build reliability over unreliable channels
- Key mechanisms: error detection, acknowledgments, sequence numbers, timers, re-transmissions
- RDT protocol evolution addresses increasingly realistic channel impairments
- Stop-and-wait protocols have poor utilization; pipelining dramatically improves performance
- Go-Back-N and Selective Repeat are two pipelined approaches with different trade-offs
- Window size constraints are critical for protocol correctness