

# TCP Congestion Control

## Study-Ready Notes

Compiled by Andrew Photinakis

October 17, 2025

## Contents

|   |          |
|---|----------|
| <b>1 TCP AIMD (Additive Increase Multiplicative Decrease)</b> | <b>2</b> |
| 1.1 Basic Approach . . . . .                                  | 2        |
| 1.2 Multiplicative Decrease Details . . . . .                 | 2        |
| 1.3 Why AIMD? . . . . .                                       | 2        |
| <b>2 TCP Congestion Control Details</b>                       | <b>3</b> |
| 2.1 Congestion Window (cwnd) . . . . .                        | 3        |
| <b>3 TCP Slow Start</b>                                       | <b>3</b> |
| 3.1 Initial Phase . . . . .                                   | 3        |
| 3.2 Exponential Growth . . . . .                              | 3        |
| <b>4 TCP Congestion Avoidance</b>                             | <b>3</b> |
| 4.1 Transition from Slow Start . . . . .                      | 3        |
| 4.2 Congestion Control States . . . . .                       | 4        |
| <b>5 TCP Congestion Control Finite State Machine</b>          | <b>4</b> |
| <b>6 TCP CUBIC</b>  | <b>4</b> |
| 6.1 Motivation for Improvement . . . . .                      | 4        |
| 6.2 CUBIC Algorithm . . . . .                                 | 5        |
| 6.3 Deployment Status . . . . .                               | 5        |
| <b>7 Bottleneck Link Concept</b>                              | <b>5</b> |
| 7.1 Network Bottlenecks . . . . .                             | 5        |
| 7.2 Key Insights . . . . .                                    | 6        |
| <b>8 Delay-Based TCP Congestion Control</b>                   | <b>6</b> |
| 8.1 Philosophy . . . . .                                      | 6        |
| 8.2 Algorithm Approach . . . . .                              | 6        |
| 8.3 Deployment Examples . . . . .                             | 6        |

|   |           |
|---|-----------|
| <b>9 Explicit Congestion Notification (ECN)</b> | <b>7</b>  |
| 9.1 Network-Assisted Approach . . . . .         | 7         |
| <b>10 TCP Fairness</b>                          | <b>7</b>  |
| 10.1 Fairness Goal . . . . .                    | 7         |
| 10.2 Fairness Analysis . . . . .                | 7         |
| 10.3 Limitations and Violations . . . . .       | 7         |
| <b>11 TCP Throughput Analysis</b>               | <b>8</b>  |
| 11.1 Average Throughput Formula . . . . .       | 8         |
| 11.2 Throughput Example . . . . .               | 8         |
| <b>12 Study Aids and Exam Preparation</b>       | <b>9</b>  |
| 12.1 Key Concepts to Master . . . . .           | 9         |
| 12.2 Practice Questions . . . . .               | 9         |
| <b>13 Summary</b>                               | <b>10</b> |

# 1 TCP AIMD (Additive Increase Multiplicative Decrease)

## 1.1 Basic Approach

- **Approach:** Senders increase sending rate until packet loss (congestion) occurs, then decrease on loss event
- **Additive Increase:** Increase sending rate by 1 MSS every RTT until loss detected
- **Multiplicative Decrease:** Cut sending rate in half at each loss event
- Creates characteristic **sawtooth pattern** - probing for available bandwidth

## 1.2 Multiplicative Decrease Details

- **TCP Reno:** Cut rate in half on loss detected by triple duplicate ACK
- **TCP Tahoe:** Cut rate to 1 MSS when loss detected by timeout

## 1.3 Why AIMD?

- Distributed, asynchronous algorithm proven to:
  - Optimize congested flow rates network-wide
  - Have desirable stability properties
  - Converge to fair allocation among competing flows

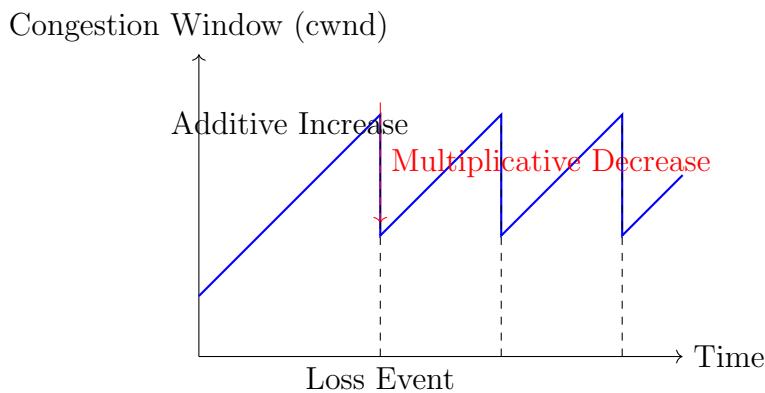


Figure 1: AIMD sawtooth behavior: probing for available bandwidth

[Summary: AIMD is TCP's core congestion control mechanism where senders gradually increase rates until congestion occurs, then dramatically reduce rates, creating a stable sawtooth pattern that fairly shares bandwidth.]

## 2 TCP Congestion Control Details

### 2.1 Congestion Window (cwnd)

- Dynamic limit on amount of unacknowledged data
- TCP sending rate approximately: Rate  $\approx \frac{\text{cwnd}}{\text{RTT}}$  bytes/sec
- Sender constraint:  $\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$

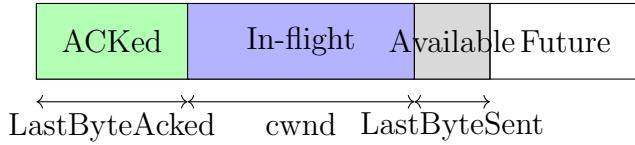


Figure 2: TCP sender sequence number space with congestion window

## 3 TCP Slow Start

### 3.1 Initial Phase

- When connection begins, increase rate exponentially until first loss
- Initial cwnd = 1 MSS
- Double cwnd every RTT
- Achieved by incrementing cwnd for every ACK received

### 3.2 Exponential Growth

$$\text{cwnd}_{\text{after RTT}} = \text{cwnd} \times 2$$

$$\text{cwnd}_{\text{after n RTTs}} = 2^n \times \text{MSS}$$

[Mnemonic: "Slow Start Speeds Up" - Despite the name, slow start actually increases the window exponentially fast, not slowly.]

## 4 TCP Congestion Avoidance

### 4.1 Transition from Slow Start

- Switch from exponential to linear increase when cwnd reaches half its pre-loss value
- **ssthresh** (slow start threshold) variable tracks this point
- On loss event:  $\text{ssthresh} = \frac{\text{cwnd}}{2}$

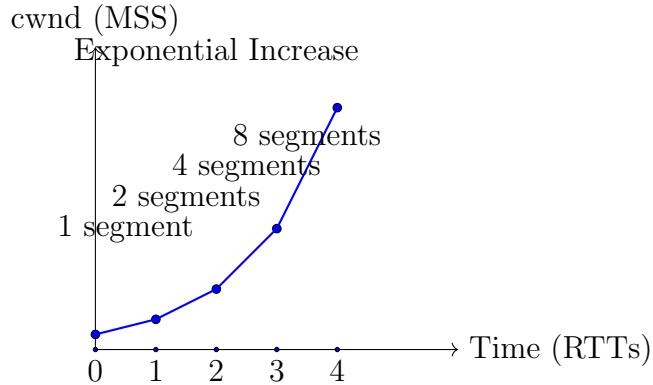


Figure 3: TCP slow start: exponential growth of congestion window

## 4.2 Congestion Control States

- **Slow Start:** Exponential increase until ssthresh
- **Congestion Avoidance:** Additive increase (AIMD)
- **Fast Recovery:** Handling duplicate ACKs

## 5 TCP Congestion Control Finite State Machine

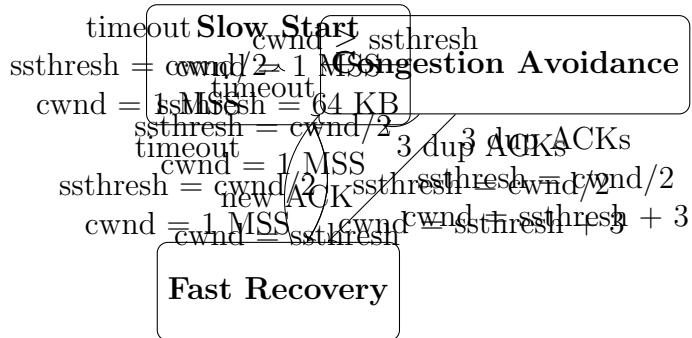


Figure 4: TCP congestion control finite state machine

## 6 TCP CUBIC

### 6.1 Motivation for Improvement

- Is there a better way than AIMD to "probe" for usable bandwidth?
- Key insight: After cutting rate on loss, initially ramp to  $W_{\max}$  faster, then approach more slowly

- $W_{\max}$ : Sending rate at which congestion loss was detected

## 6.2 CUBIC Algorithm

- $K$ : Point in time when window will reach  $W_{\max}$
- Increase window as cube of distance from current time to  $K$
- Larger increases when further from  $K$ , smaller increases when nearer

$$W(t) = C(t - K)^3 + W_{\max}$$

Where  $C$  is a scaling constant.

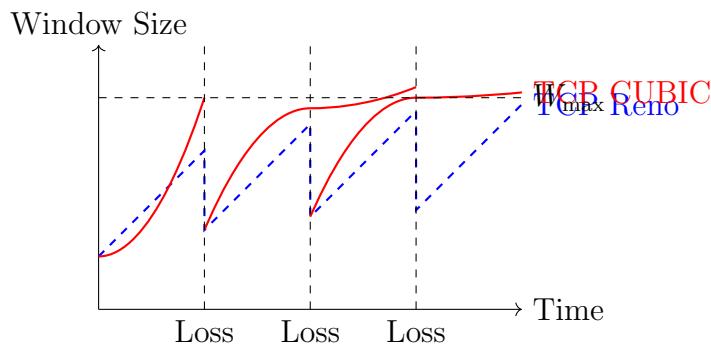


Figure 5: TCP CUBIC vs classic TCP: CUBIC achieves higher throughput

## 6.3 Deployment Status

- Default TCP in Linux
- Most popular TCP for popular web servers
- Provides better performance on high-bandwidth, high-latency networks

[Summary: TCP CUBIC improves upon AIMD by using a cubic growth function that initially recovers quickly after loss then approaches the previous maximum more cautiously, providing better performance on modern networks.]

# 7 Bottleneck Link Concept

## 7.1 Network Bottlenecks

- TCP increases sending rate until packet loss occurs at some router's output
- This congested router is the **bottleneck link**
- Understanding congestion requires focusing on bottleneck links

## 7.2 Key Insights

- Increasing TCP sending rate beyond bottleneck capacity won't increase end-to-end throughput
- Increasing TCP sending rate will increase measured RTT due to queueing
- Goal: "Keep the end-end pipe just full, but not fuller"

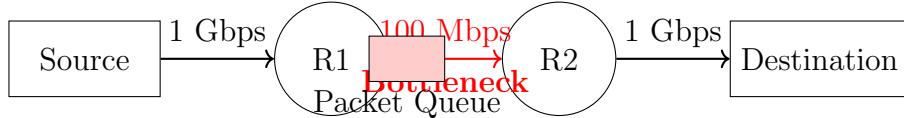


Figure 6: Bottleneck link concept: the slowest link determines maximum throughput

## 8 Delay-Based TCP Congestion Control

### 8.1 Philosophy

- Keep sender-to-receiver pipe "just full enough, but no fuller"
- Keep bottleneck link busy transmitting, but avoid high delays/buffering
- Congestion control without inducing/forcing loss

### 8.2 Algorithm Approach

- Measure  $RTT_{min}$ : Minimum observed RTT (uncongested path)
- Calculate uncongested throughput:  $\frac{cwnd}{RTT_{min}}$
- Compare measured throughput with uncongested throughput:
  - If "very close": Increase cwnd linearly (path not congested)
  - If "far below": Decrease cwnd linearly (path congested)

### 8.3 Deployment Examples

- BBR (Bottleneck Bandwidth and Round-trip propagation time)
- Deployed on Google's internal backbone network
- Better performance for latency-sensitive applications

## 9 Explicit Congestion Notification (ECN)

### 9.1 Network-Assisted Approach

- Routers provide direct congestion feedback
- Two bits in IP header (ToS field) marked by network router
- Congestion indication carried to destination
- Destination sets ECE bit on ACK to notify sender

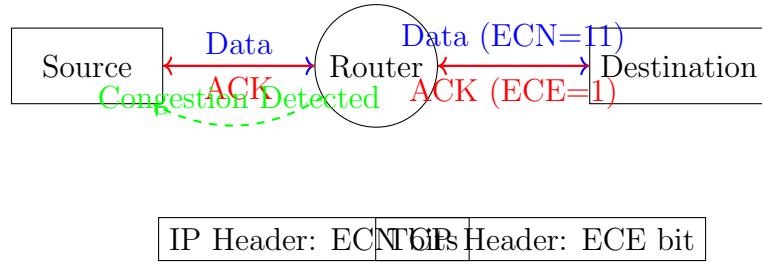


Figure 7: Explicit Congestion Notification (ECN) mechanism

## 10 TCP Fairness

### 10.1 Fairness Goal

- If  $K$  TCP sessions share bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$
- TCP converges to fair allocation under idealized conditions

### 10.2 Fairness Analysis

- Additive increase gives slope of 1 as throughput increases
- Multiplicative decrease reduces throughput proportionally
- Result: Competing flows converge to equal shares

### 10.3 Limitations and Violations

- **UDP applications:** Multimedia apps don't use TCP congestion control
- **Parallel TCP connections:** Applications can open multiple connections
- Example: Link with rate  $R$  and 9 existing connections:

- New app with 1 TCP gets rate  $R/10$
- New app with 11 TCPs gets rate  $R/2$
- No "Internet police" enforcing congestion control use

| Congestion Control Type | Mechanism             | Characteristics                   |
|-------------------------|-----------------------|-----------------------------------|
| Classic TCP (Reno)      | Loss-based AIMD       | Sawtooth pattern, proven fairness |
| TCP CUBIC               | Cubic growth function | Better high-speed performance     |
| Delay-based (BBR)       | RTT measurements      | Lower latency, no forced loss     |
| ECN-enabled             | Network feedback      | Early congestion notification     |

Table 1: Comparison of TCP congestion control variants

## 11 TCP Throughput Analysis

### 11.1 Average Throughput Formula

Ignoring slow start and assuming always data to send:

$$\text{Average TCP throughput} = \frac{3}{4} \cdot \frac{W}{\text{RTT}} \text{ bytes/sec}$$

Where:

- $W$ : Window size where loss occurs
- Average window size is  $\frac{3}{4}W$
- Average throughput is  $\frac{3}{4}W$  per RTT

### 11.2 Throughput Example

For a connection with:

- Loss window  $W = 10$  MSS
- RTT = 100 ms
- MSS = 1460 bytes

$$\text{Average throughput} = \frac{3}{4} \cdot \frac{10 \times 1460}{0.1} = 109,500 \text{ bytes/sec} \approx 876 \text{ kbps}$$

[Concept Map: TCP Congestion Control → AIMD (core algorithm) + Slow Start (initial phase) + Congestion Avoidance (steady state) + Enhancements (CUBIC, BBR, ECN) → Goals: Efficiency + Fairness + Stability → Prevents congestion collapse while maximizing network utilization.]

## 12 Study Aids and Exam Preparation

### 12.1 Key Concepts to Master

- Understand AIMD mechanism and why it creates sawtooth pattern
- Differentiate between slow start and congestion avoidance
- Explain TCP CUBIC improvements over classic TCP
- Compare loss-based vs delay-based congestion control
- Describe ECN mechanism and benefits
- Analyze TCP fairness and its limitations

### 12.2 Practice Questions

1. Explain the **AIMD mechanism** and draw the characteristic sawtooth pattern. Why does this pattern emerge, and what are its stability properties?
2. Compare **TCP slow start** and **congestion avoidance**. When does the transition occur, and what triggers it?
3. Calculate the **average TCP throughput** for a connection that experiences loss when cwnd = 16 MSS, with RTT = 50 ms and MSS = 1500 bytes.
4. Describe how **TCP CUBIC** improves upon classic TCP congestion control. What is the key insight behind its cubic growth function?
5. Explain the **fairness limitations** of TCP congestion control. How can applications "cheat" the system, and what are the implications?

[Mnemonic: "AIMD: Add Incrementally, Multiply Down" - Add 1 MSS per RTT when increasing, multiply by 0.5 when decreasing.]

## 13 Summary

- TCP congestion control prevents network collapse while maximizing utilization
- AIMD (Additive Increase Multiplicative Decrease) is the core algorithm
- Slow start provides exponential initial growth
- Congestion avoidance provides linear increase after threshold
- TCP CUBIC improves performance on high-speed networks
- Delay-based approaches (BBR) avoid forced packet loss
- ECN enables network-assisted congestion notification
- TCP achieves fairness under idealized conditions
- Various enhancements address limitations of classic approaches
- Throughput can be modeled mathematically based on loss window and RTT