

# Message Passing Interface (MPI)

## Study-Ready Notes

Compiled by Andrew Photinakis

## Contents

<b>1</b>	<b>Introduction to MPI</b>	<b>2</b>
1.1	What is MPI? . . . . .	2
1.2	Parallel Computing Systems . . . . .	2
1.3	MPI Program Structure . . . . .	2
<b>2</b>	<b>MPI Communication Fundamentals</b>	<b>2</b>
2.1	Basic Send and Receive Operations . . . . .	2
2.2	Communicators and Groups . . . . .	3
<b>3</b>	<b>Point-to-Point Communication</b>	<b>3</b>
3.1	Communication Types . . . . .	3
3.2	System Buffers . . . . .	3
3.3	Blocking vs Non-blocking Operations . . . . .	4
3.3.1	Blocking Operations . . . . .	4
3.3.2	Non-blocking Operations . . . . .	4
<b>4</b>	<b>Collective Communication</b>	<b>4</b>
4.1	Types of Collective Operations . . . . .	4
4.2	Important Collective Routines . . . . .	4
4.2.1	Broadcast (MPI_Bcast) . . . . .	4
4.2.2	Scatter (MPI_Scatter) . . . . .	5
4.2.3	Gather (MPI_Gather) . . . . .	5
4.2.4	Reduce (MPI_Reduce) . . . . .	5
4.2.5	Allreduce (MPI_Allreduce) . . . . .	5
<b>5</b>	<b>Derived Data Types</b>	<b>5</b>
5.1	Continuous Derived Data Type . . . . .	5
5.2	Vector Derived Data Type . . . . .	5
5.3	Indexed Derived Data Type . . . . .	6
5.4	Struct Derived Data Type . . . . .	6

<b>6 Groups and Communicators</b>	<b>6</b>
6.1 Definitions . . . . .	6
6.2 Group Management Operations . . . . .	6
6.3 Communicator Splitting . . . . .	7
<b>7 Virtual Topologies</b>	<b>7</b>
7.1 Concept and Purpose . . . . .	7
7.2 Benefits . . . . .	7
7.3 Cartesian Topology Example . . . . .	7
<b>8 MPI Programming Examples</b>	<b>8</b>
8.1 Monte Carlo Pi Calculation . . . . .	8
8.1.1 Mathematical Basis . . . . .	8
8.1.2 Parallel Implementation Strategy . . . . .	8
8.2 Nearest Neighbor Exchange . . . . .	8
<b>9 MPI Advantages and Characteristics</b>	<b>8</b>
9.1 Key Features . . . . .	8
<b>10 Study Aids</b>	<b>9</b>

# 1 Introduction to MPI

## 1.1 What is MPI?

- Message Passing Interface (MPI) - a specification, not a library
- Message-passing parallel programming model: data moved from address space of one process to another through cooperative operations
- Well-known implementations: Open MPI and MVAPICH2
- All parallelism is explicit: programmer responsible for identifying parallelism and implementing parallel algorithms

[Summary: MPI is a standardized message-passing specification for parallel programming where processes communicate by explicitly sending and receiving messages between their separate address spaces.]

## 1.2 Parallel Computing Systems

- Distributed Memory: Computers in a network
- Distributed Shared Memory: Computers in a cluster
- Multiprocessor systems
- Multicore systems
- In-cloud and Edge computing

## 1.3 MPI Program Structure

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int numtasks, rank, dest, source, rc, count, tag=1;
    char inmsg, outmsg='x';
    MPI_Status Stat;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if (rank == 0) {
        dest = 1; source = 1;
        rc = MPI_Send(&outmsg, 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
```

```

        rc = MPI_Recv(&inmsg, 1, MPI_CHAR, source, tag, MPI_COMM_WORLD, &Stat);
    }
    else if (rank == 1) {
        dest = 0; source = 0;
        rc = MPI_Recv(&inmsg, 1, MPI_CHAR, source, tag, MPI_COMM_WORLD, &Stat);
        rc = MPI_Send(&outmsg, 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
    }

    MPI_Get_count(&Stat, MPI_CHAR, &count);
    printf("Task %d: Received %d char(s) from task %d with tag %d\n",
           rank, count, Stat.MPI_SOURCE, Stat.MPI_TAG);

    MPI_Finalize();
}

```

[Summary: Basic MPI programs initialize with MPI\_Init, get process count and rank, perform communication operations, and finalize with MPI\_Finalize. Processes are identified by unique ranks within communicators.]

## 2 MPI Communication Fundamentals

### 2.1 Basic Send and Receive Operations

- **MPI\_Send(void \*sendbuf, int nelems, int dest, int tag, MPI\_Comm comm)**
- **MPI\_Recv(void \*recvbuf, int nelems, int source, int tag, MPI\_Comm comm, MPI\_Status \*status)**
- Parameters:
  - sendbuf/recvbuf: pointers to data buffers
  - nelems: number of data elements
  - dest/source: process identifiers
  - tag: message identifier (non-negative integer)
  - comm: communicator (usually MPI\_COMM\_WORLD)

### 2.2 Communicators and Groups

- **Communicators:** Define which processes can communicate
- **MPI\_COMM\_WORLD:** Predefined communicator including all processes
- **Rank:** Unique integer identifier for each process within a communicator
- Ranks are contiguous starting from 0

- Used to specify source/destination of messages and control program flow

[Summary: MPI uses communicators to define communication groups, with ranks identifying individual processes for message routing and conditional execution.]

## 3 Point-to-Point Communication

### 3.1 Communication Types

- Synchronous send
- Blocking send/blocking receive
- Non-blocking send/non-blocking receive
- Buffered send
- Combined send/receive
- "Ready" send
- Any send type can pair with any receive type

### 3.2 System Buffers

- Buffer space reserved for data in transit
- Managed entirely by MPI library (opaque to programmer)
- Finite resource that can be exhausted
- Can exist on sending side, receiving side, or both
- Allows send-receive operations to be asynchronous

[Concept Map: Point-to-Point Communication → Blocking/Non-blocking → Synchronous/Asynchronous → Buffered/Non-buffered → System Buffer Management]

### 3.3 Blocking vs Non-blocking Operations

#### 3.3.1 Blocking Operations

- **Blocking Send:** Returns only when safe to modify application buffer
- **Blocking Receive:** Returns only when data has arrived and is ready
- Can be synchronous (handshaking) or asynchronous (using system buffers)

### 3.3.2 Non-blocking Operations

- Return immediately without waiting for communication events
- Simply "request" MPI to perform operation when able
- Unsafe to modify application buffer until operation completes
- Use "wait" routines (MPI\_Wait, MPI\_Test) to check completion
- Used to overlap computation with communication

[Mnemonic: Blocking = Wait for completion, Non-blocking = Fire and (maybe) forget until checked]

## 4 Collective Communication

### 4.1 Types of Collective Operations

- **Synchronization:** Processes wait until all reach synchronization point (MPI\_Barrier, MPI\_Waitall)
- **Data Movement:** Broadcast, scatter/gather, all to all
- **Collective Computation (Reductions):** One member collects data and performs operations (min, max, add, multiply)

### 4.2 Important Collective Routines

#### 4.2.1 Broadcast (MPI\_Bcast)

- Broadcasts message from one task to all others in communicator
- `MPI_Bcast(&buffer, count, datatype, root, comm)`

#### 4.2.2 Scatter (MPI\_Scatter)

- Sends chunks of data from one task to all tasks
- `MPI_Scatter (sendbuf, sendcnt, sendtype, recvbuf, recvcnt, recvtype, root, comm)`

#### 4.2.3 Gather (MPI\_Gather)

- Gathers data from all tasks to a single task
- Inverse of scatter operation

#### 4.2.4 Reduce (MPI\_Reduce)

- Performs reduction (sum, max, min, etc.) across all tasks
- Result stored in one task
- `MPI_Reduce(sendbuf, recvbuf, count, datatype, op, root, comm)`

#### 4.2.5 Allreduce (MPI\_Allreduce)

- Performs reduction and stores result across all tasks

[Summary: Collective operations involve all processes in a communicator for synchronization, data movement, or collective computation, with routines like broadcast, scatter/gather, and reductions.]

### 5 Derived Data Types

#### 5.1 Continuous Derived Data Type

- Represents contiguous blocks of data
- `MPI_Type_contiguous(count, oldtype, &newtype)`
- Example: Representing a row of a 2D array

#### 5.2 Vector Derived Data Type

- Represents regularly spaced data blocks
- `MPI_Type_vector(count, blocklength, stride, oldtype, &newtype)`
- Example: Representing a column of a 2D array

#### 5.3 Indexed Derived Data Type

- Represents irregular data patterns
- `MPI_Type_indexed(count, array_of_blocklengths, array_of_displacements, oldtype, &newtype)`

#### 5.4 Struct Derived Data Type

- Represents heterogeneous data structures
- `MPI_Type_struct(count, array_of_blocklengths, array_of_offsets, array_of_types, &newtype)`
- Example: Representing a particle with position, velocity, and type information

[Mnemonic: Continuous = Contiguous blocks, Vector = Regular spacing, Indexed = Irregular patterns, Struct = Mixed data types]

## 6 Groups and Communicators

### 6.1 Definitions

- **Group:** Ordered set of processes with unique integer ranks (0 to N-1)
- **Communicator:** Encompasses a group of processes that can communicate
- **MPI\_COMM\_WORLD:** Default communicator containing all processes
- Groups and communicators are dynamic - can be created/destroyed during execution
- Processes can be in multiple groups/communicators with different ranks

### 6.2 Group Management Operations

1. Extract global group: `MPI_Comm_group(MPI_COMM_WORLD, &group)`
2. Create new group: `MPI_Group_incl(group, n, ranks, &newgroup)`
3. Create new communicator: `MPI_Comm_create(comm, group, &newcomm)`
4. Get new rank: `MPI_Comm_rank(newcomm, &rank)`
5. Free resources: `MPI_Comm_free()`, `MPI_Group_free()`

### 6.3 Communicator Splitting

- `MPI_Comm_split(comm, color, key, &newcomm)`
- Groups processes by color, sorts by key within each color group
- Useful for organizing processes by function or data partitioning

[Summary: Groups define process sets, communicators enable communication within those sets, and MPI provides operations to dynamically create and manage these entities for flexible parallel programming.]

## 7 Virtual Topologies

### 7.1 Concept and Purpose

- Virtual topology: Mapping/ordering of MPI processes into geometric shapes
- Types: Cartesian (grid) and Graph topologies

- Virtual - no relation to physical machine structure required
- Must be programmed by application developer

## 7.2 Benefits

- **Convenience:** Match communication patterns to application needs
- **Efficiency:** Potential optimization for specific hardware architectures
- Example: Cartesian topology for grid-based computations with nearest-neighbor communication

## 7.3 Cartesian Topology Example

- Processes arranged in 2D grid with row-major ordering
- Each process identified by (row, column) coordinates
- Supports operations like `MPI_Cart_sub()` for creating sub-grids

[Concept Map: Virtual Topologies → Cartesian/Graph Types → Process Mapping → Communication Patterns → Potential Hardware Optimization]

# 8 MPI Programming Examples

## 8.1 Monte Carlo Pi Calculation

### 8.1.1 Mathematical Basis

- Circle area:  $\pi r^2$ , Square area:  $4r^2$
- Ratio:  $\frac{\pi r^2}{4r^2} = \frac{\pi}{4}$
- Random points:  $M \approx \frac{N\pi}{4}$
- Approximation:  $\pi \approx \frac{4M}{N}$

### 8.1.2 Parallel Implementation Strategy

- Divide total points among available tasks
- Each task calculates points in circle for its subset
- Master task collects results using send/receive operations
- No data dependencies between tasks

## 8.2 Nearest Neighbor Exchange

```
// Non-blocking ring communication
MPI_Irecv(&buf[0], 1, MPI_INT, prev, tag1, MPI_COMM_WORLD, &reqs[0]);
MPI_Irecv(&buf[1], 1, MPI_INT, next, tag2, MPI_COMM_WORLD, &reqs[1]);
MPI_Isend(&rank, 1, MPI_INT, prev, tag2, MPI_COMM_WORLD, &reqs[2]);
MPI_Isend(&rank, 1, MPI_INT, next, tag1, MPI_COMM_WORLD, &reqs[3]);

// Do work while communications progress
MPI_Waitall(4, reqs, stats);
```

[Summary: MPI enables various parallel algorithms through point-to-point and collective communication, with examples including Monte Carlo simulations and nearest-neighbor exchanges in ring topologies.]

# 9 MPI Advantages and Characteristics

## 9.1 Key Features

- **Practical:** Easy to implement on any system
- **Portable:** Works on networks, clusters, multiprocessors
- **Efficient:** Programmer has more control over communication
- **Flexible:** Easy to scale problem and processor size
- **Explicit Parallelism:** Burden on programmer to identify and manage parallelism

# 10 Study Aids

## Key Concepts to Master

- Difference between blocking and non-blocking operations
- Purpose and usage of communicators and groups
- Collective communication operations and when to use them
- Derived data types and their applications
- Virtual topologies and their benefits
- MPI program structure and initialization/finalization

## Common MPI Functions Reference

- Environment: MPI\_Init, MPI\_Finalize, MPI\_Comm\_size, MPI\_Comm\_rank
- Point-to-point: MPI\_Send, MPI\_Recv, MPI\_Isend, MPI\_Irecv, MPI\_Wait
- Collective: MPI\_Bcast, MPI\_Scatter, MPI\_Gather, MPI\_Reduce, MPI\_Barrier
- Groups/Communicators: MPI\_Comm\_split, MPI\_Comm\_create, MPI\_Group\_incl

[Mnemonic: MPI Basics - Init, Size, Rank, Send, Recv, Finalize. Collective Ops - Bcast, Scatter, Gather, Reduce. Advanced - Groups, Topologies, Derived Types]