

Graph Algorithms and Parallel Computing

Study-Ready Notes

Compiled by Andrew Photinakis

October 2nd, 2025

Contents

1	Graph Representations	3
1.1	Undirected Graph Representation	3
1.2	Directed Graph Representation	3
2	Graph Traversal Algorithms	3
2.1	Depth-First Search (DFS)	3
2.2	Breadth-First Search (BFS)	4
3	Minimum-Cost Spanning Trees (MCST)	4
3.1	Definition and Applications	4
3.2	Prim's Algorithm	5
4	Shortest Path Algorithms	5
4.1	Single-Source Shortest Paths (Dijkstra's Algorithm)	5
4.2	All-Pairs Shortest Paths	6
5	Transitive Closure	6
5.1	Definition and Applications	6
5.2	Warshall's Algorithm	6
6	Matrix Operations on Graphs	7
6.1	Connectivity and Path Counting	7
6.2	All-Pairs Shortest Paths via Matrix Operations	7
7	Optimal Binary Search Trees (OBST)	7
7.1	Problem Definition	7
7.2	Recursive Structure	8
7.3	Bottom-Up Computation	8
7.4	Parallel OBST Computation	8

8	Subgraph Matching	8
8.1	Problem Definition	8
8.2	Query Decomposition	9
9	Process Assignment and Scheduling	9
9.1	Basic Concepts	9
9.2	Critical Factors	9
9.3	System Characteristics	10
10	Decomposition Strategies	10
10.1	Domain Decomposition	10
10.2	Functional Decomposition	10
11	Load Balancing	11
11.1	Static Load Balancing	11
11.2	Dynamic Load Balancing	11
11.3	Process Migration	11

1 Graph Representations

1.1 Undirected Graph Representation

- **Adjacency Matrix:** Square matrix where entry $(i,j) = 1$ if vertices i and j are connected, 0 otherwise
- **Adjacency List:** For each vertex, list of adjacent vertices

Example G1 (Undirected Graph)

Adjacency Matrix:

$$\begin{bmatrix} 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \end{bmatrix}$$

Figure 1: Undirected graph with vertices A,B,C,D,E

1.2 Directed Graph Representation

Example G2 (Directed Graph)

Adjacency Matrix:

$$\begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 & 0 \end{bmatrix}$$

Figure 2: Directed graph with vertices P,Q,R,S,T

[Summary] Graph representations include adjacency matrices (good for dense graphs) and adjacency lists (good for sparse graphs). Directed graphs have asymmetric matrices.

2 Graph Traversal Algorithms

2.1 Depth-First Search (DFS)

- Explores as far as possible along each branch before backtracking

- Uses stack (implicit or explicit) for traversal
- Applications: Cycle detection, topological sorting, maze solving

Example DFS Tree:

- $A \rightarrow D, E, B$
- $B \rightarrow A, E, C$
- $C \rightarrow D, E, B$
- $D \rightarrow A, C$
- $E \rightarrow A, B, C$

2.2 Breadth-First Search (BFS)

- Explores all neighbors at current depth before moving deeper
- Uses queue for traversal
- Applications: Shortest path in unweighted graphs, social networks

Example BFS Tree:

- $A \rightarrow B, E, D$
- $B \rightarrow E, C$
- $E \rightarrow C$
- $D \rightarrow C$

[Summary] DFS goes deep first using stack, BFS goes wide first using queue. DFS finds paths, BFS finds shortest paths in unweighted graphs.

3 Minimum-Cost Spanning Trees (MCST)

3.1 Definition and Applications

- **Spanning Tree:** Connected subgraph containing all vertices with no cycles
- **Minimum-Cost Spanning Tree:** Spanning tree with minimum total edge weight
- **Applications:** Network design, circuit wiring, clustering

Network Example:

- Computer network with bidirectional links
- Each link has positive cost (message sending cost)
- Broadcast message from arbitrary computer
- Goal: Minimize total broadcast cost

3.2 Prim's Algorithm

```
def prim_mst(graph, start_node):
    mst = set()
    visited = {start_node}
    edges = [
        (cost, start_node, to)
        for to, cost in graph[start_node].items()
    ]
    heapify(edges)

    while edges and len(visited) < len(graph):
        cost, frm, to = heappop(edges)
        if to not in visited:
            visited.add(to)
            mst.add((frm, to, cost))
            for to_next, cost2 in graph[to].items():
                if to_next not in visited:
                    heappush(edges, (cost2, to, to_next))
    return mst
```

Algorithm Steps:

1. Start with any node as root
2. Grow tree greedily by adding cheapest edge connecting tree to outside vertex
3. Repeat until all vertices are included

Complexity: $O(E \log V)$ with binary heap

[Summary] MCST finds minimum weight tree spanning all vertices. Prim's algorithm grows tree greedily from start node.

4 Shortest Path Algorithms

4.1 Single-Source Shortest Paths (Dijkstra's Algorithm)

- Finds shortest paths from source vertex to all other vertices
- Works for weighted graphs with non-negative weights
- Based on greedy principle

Algorithm:

1. Initialize $d[v] = 0$ for source, ∞ for others
2. For each vertex, compute: $d[x] = \min\{d[x], d[v] + w(v, x)\}$

3. Always pick vertex with minimum distance

Mathematical Formulation:

$$d[x] = \min\{d[x], d[v] + w(v, x)\}, \text{ where } v, x \in V$$

4.2 All-Pairs Shortest Paths

Recursive Solution:

$$dist(i, j) = \begin{cases} w(i, j) & \text{if } k = 0 \\ \min\{dist(i, j), [dist(i, k) + dist(k, j)]\} & \text{if } k \geq 1 \end{cases}$$

Matrix Operations Approach:

- Replace 'multiply' by 'ADD'
- Replace 'add' by 'MINIMUM'
- Ignore infinity entries

[Summary] Dijkstra finds single-source shortest paths, all-pairs uses dynamic programming. Both use greedy/minimization principles.

5 Transitive Closure

5.1 Definition and Applications

- **Transitive Closure:** Directed graph where edge (i,j) exists if there's a directed path from i to j in original graph
- **Security Application:** Identify all users with permission (direct or indirect) to access accounts
- Many applications in database systems, compiler optimization

5.2 Warshall's Algorithm

procedure WARSHALL($G=[V, E]$)

Input: $n \times n$ matrix A representing adjacency

Output: transitive closure matrix T

```

for i ← 1 to n do
  for j ← 1 to n do
    t[i, j] ← a(i, j)

```

```

for k ← 1 to n do
  for i ← 1 to n do

```

```

    for j ← 1 to n do
        if NOT t[i,j] then
            t[i,j] ← t[i,k] AND t[k,j]
    return T

```

Complexity: $\Theta(n^3)$

Improvement: Algorithm can be optimized for better performance

[Summary] Transitive closure identifies all reachable pairs in a graph. Warshall's algorithm computes it in cubic time using dynamic programming.

6 Matrix Operations on Graphs

6.1 Connectivity and Path Counting

Paths of Length 2:

- Replace 'multiply' by 'AND' and 'add' by 'OR' for existence
- Keep 'add' and replace 'multiply' by 'AND' for counting

Example:

$$C^2 = C \times C = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \end{bmatrix} \times \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

6.2 All-Pairs Shortest Paths via Matrix Operations

Operations:

- Replace 'multiply' by 'ADD'
- Replace 'add' by 'MINIMUM'
- Handle infinity entries appropriately

[Summary] Matrix operations can compute connectivity and shortest paths by redefining multiplication and addition operations.

7 Optimal Binary Search Trees (OBST)

7.1 Problem Definition

- Given keys with access probabilities
- Find BST arrangement that minimizes expected access cost
- Cost = $\sum (\text{probability} \times \text{depth})$ for all keys

Example: Keys A,B,C,D with probabilities (0.1, 0.2, 0.4, 0.3)

7.2 Recursive Structure

$$c[i, j] = c[i, k - 1] + c[k + 1, j] + \sum_{s=i}^j p_s$$

Where:

- $c[i, j]$ = cost of optimal BST for keys i through j
- k = root of subtree
- $\sum p_s$ = sum of probabilities in current subtree

7.3 Bottom-Up Computation

Base Cases:

- $c[i, j] = 0$ if $i = 0$ or $i \geq j$
- $c[i, j] = p_i$ if $i = j$

Example Computation with P=0.2, Q=0.4, R=0.1, S=0.3:

- $C[1,1] = 0.2$, $C[2,2] = 0.4$, $C[3,3] = 0.1$, $C[4,4] = 0.3$
- $C[1,2] = 0.8$ (Q root), $C[2,3] = 0.6$ (Q root), $C[3,4] = 0.5$ (S root)

7.4 Parallel OBST Computation

- Compute diagonals in parallel
- $C(1,2)$, $C(2,3)$, $C(3,4)$ on $n-1$ processors
- $C(1,3)$, $C(2,4)$ on $n-2$ processors
- Load balancing needed for initial unbalanced assignments

[Summary] OBST minimizes expected search cost using dynamic programming. Parallel computation processes matrix diagonals concurrently.

8 Subgraph Matching

8.1 Problem Definition

- Given data graph G and query graph Q
- Find all subgraphs of G isomorphic to Q
- Applications: Social networks, web graphs, relational databases

Formal Definition:

- $G(V, E)$, $Q(V_q, E_q)$
- Find subgraph $g(V_g, E_g)$ where $V_q \rightarrow V_g$ and $E_q \rightarrow E_g$

8.2 Query Decomposition

- Decompose complex query into simpler components (twigs)
- Each processor searches for specific twig in distributed graph
- Handle large graphs: $|E| = O(10^9)$ and $|V| = O(10^8)$

Parallelization Strategy:

- Distribute G across computers
- Each computer searches for assigned twig pattern
- Combine results from all processors

[Summary] Subgraph matching finds pattern occurrences in large graphs. Parallel approach decomposes query and distributes search.

9 Process Assignment and Scheduling

9.1 Basic Concepts

- **Assignment:** Processes to processing elements (WHERE)
- **Scheduling:** When to execute each task (WHEN)
- **Programming Models:** SPMD, MPMD, Shared Memory, Message Passing

9.2 Critical Factors

Granularity:

- Coarse vs Fine grain
- Ratio of computation to communication
- Higher ratio \rightarrow better speedup and efficiency

Overheads:

- Coordination costs
- Synchronization
- Data communication

Scalability:

- Proportionate speedup with more processors
- Affected by memory-CPU bandwidth, network, algorithm characteristics

9.3 System Characteristics

Processor Types:

- **Homogeneous:** Identical processors, uniform costs
- **Heterogeneous:** Varying capabilities, speeds, resources

Network Types:

- Homogeneous/heterogeneous communication bandwidth
- Mobile systems with disconnections

Total Cost Calculation:

$$\text{Total Cost} = \text{computing costs} + \text{communication costs}$$

[Summary] Process assignment and scheduling consider granularity, overheads, scalability. Systems can be homogeneous or heterogeneous.

10 Decomposition Strategies

10.1 Domain Decomposition

- Divide data into discrete chunks
- Each process works on portion of data
- Examples: Matrix operations, image processing
- Maintain high computation/communication ratio (R/C)

10.2 Functional Decomposition

- Each processor performs different function
- Examples: Signal processing pipelines
- Match system (R, C) to application (r, c) characteristics

[Summary] Domain decomposition divides data, functional decomposition divides tasks. Both aim to optimize computation/communication ratio.

11 Load Balancing

11.1 Static Load Balancing

- Fixed policy based on a priori knowledge
- Does not adjust to system state changes

Advantages:

- Simple, low cost
- Easy session state management

Disadvantages:

- Cannot adapt to dynamic changes
- May lead to poor resource utilization

11.2 Dynamic Load Balancing

- Adjusts based on current system state
- Handles uncertainty in execution times, resource availability

Types:

- **Sender-initiated:** Overloaded nodes send work
- **Receiver-initiated:** Underloaded nodes request work
- **Centralized vs Decentralized**

Challenges:

- State monitoring overhead
- Network topology considerations
- Process migration costs

11.3 Process Migration

Migration Process:

1. Migration request initiated
2. Process suspended on source host
3. Process state transferred to destination

4. Process resumes execution on new host

Migration Scenarios:

- Host A has no more work → finds work elsewhere
- Host cannot proceed due to resource constraints

[Summary] Static balancing is simple but inflexible. Dynamic balancing adapts but has overheads. Process migration enables load redistribution.

Exam Questions

1. Compare and contrast adjacency matrix vs adjacency list representations. When would you use each?
2. Explain the greedy principle in Prim's algorithm and Dijkstra's algorithm with examples.
3. Derive the recursive formula for optimal BST cost and explain each term.
4. Compare static vs dynamic load balancing in terms of overhead, adaptability, and suitable applications.
5. Describe the process migration mechanism and discuss two scenarios where it would be beneficial.

[Mnemonic]

- **DFS**: Deep First Search (Stack)
- **BFS**: Broad First Search (Queue)
- **MCST**: Minimum Cost Spanning Tree
- **OBST**: Optimal Binary Search Tree

[Concept Map]

- Graph Algorithms → Representations → Traversal → Shortest Paths → Spanning Trees
- Scheduling → Assignment → Load Balancing → Static/Dynamic → Migration
- Parallel Computing → Decomposition → Domain/Functional → Granularity → Scalability